

Interpreting *cis*-regulatory mechanisms from genomic deep neural networks using surrogate models

Received: 27 November 2023

Evan E. Seitz , David M. McCandlish , Justin B. Kinney  & Peter K. Koo  

Accepted: 9 May 2024

Published online: 21 June 2024

 Check for updates

Deep neural networks (DNNs) have greatly advanced the ability to predict genome function from sequence. However, elucidating underlying biological mechanisms from genomic DNNs remains challenging. Existing interpretability methods, such as attribution maps, have their origins in non-biological machine learning applications and therefore have the potential to be improved by incorporating domain-specific interpretation strategies. Here we introduce SQUID (Surrogate Quantitative Interpretability for Deepnets), a genomic DNN interpretability framework based on domain-specific surrogate modelling. SQUID approximates genomic DNNs in user-specified regions of sequence space using surrogate models—simpler quantitative models that have inherently interpretable mathematical forms. SQUID leverages domain knowledge to model *cis*-regulatory mechanisms in genomic DNNs, in particular by removing the confounding effects that nonlinearities and heteroscedastic noise in functional genomics data can have on model interpretation. Benchmarking analysis on multiple genomic DNNs shows that SQUID, when compared to established interpretability methods, identifies motifs that are more consistent across genomic loci and yields improved single-nucleotide variant-effect predictions. SQUID also supports surrogate models that quantify epistatic interactions within and between *cis*-regulatory elements, as well as global explanations of *cis*-regulatory mechanisms across sequence contexts. SQUID thus advances the ability to mechanistically interpret genomic DNNs.

Deep neural networks (DNNs) are increasingly being used to analyse and biologically interpret functional genomics data. DNNs have demonstrated remarkable success at predicting diverse genomic activities from primary genome sequences, including mRNA expression levels¹, transcription start site activity^{2,3}, mRNA splicing patterns⁴, protein-DNA binding⁵, chromatin accessibility⁶ and chromatin conformation⁷. It is widely believed that the success of genomic DNNs reflects their ability to accurately model complex biological mechanisms.

However, interpreting genomic DNNs in terms of biological mechanisms remains difficult.

A variety of post hoc attribution methods have been developed for aiding the interpretation of genomic DNNs^{8,9}. The most commonly used attribution methods produce attribution maps, which quantify the position-specific effects that variant nucleotides in a sequence of interest have on DNN predictions. Attribution maps are often mechanistically interpreted by identifying motifs, that is, recurrent sequence

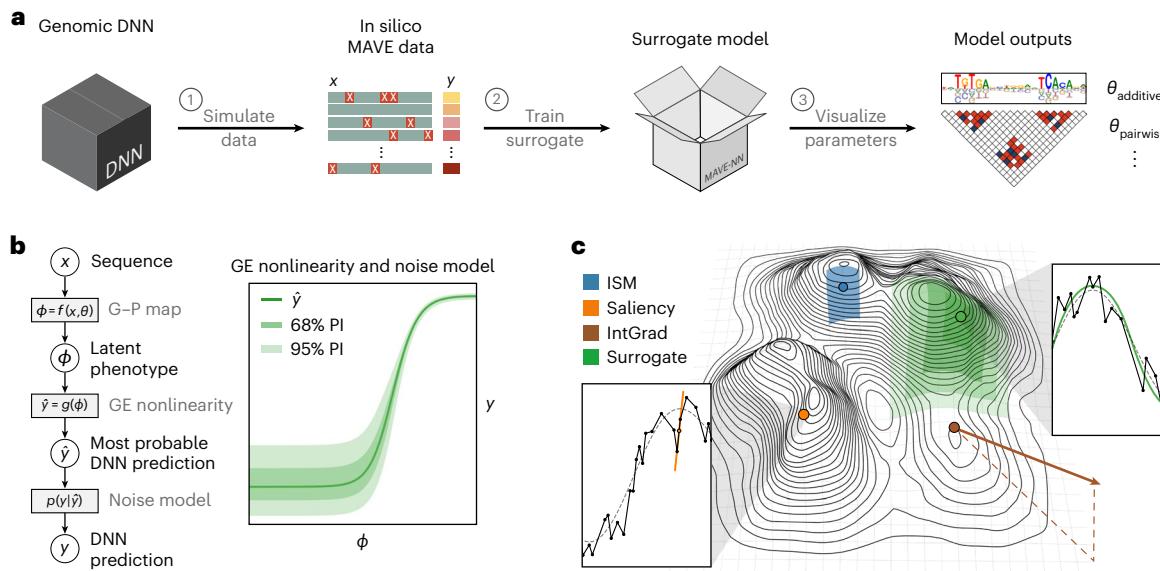


Fig. 1 | Overview of SQUID. **a**, Schematic of the SQUID modelling framework. Analysis using SQUID comprises three main steps: (1) generate an in silico MAVE dataset; (2) train a surrogate model on the MAVE dataset; and (3) visualize parameters of the surrogate model to uncover biological mechanisms. **b**, Structure of the latent phenotype surrogate models supported by SQUID. PI, prediction interval. **c**, Schematic diagram of a DNN function on a two-dimensional projection of sequence space. Each point in the plane corresponds to a unique sequence, and elevations represent DNN predictions. Green region

schematizes the ability of surrogate models to approximate the DNN function over an extended region of sequence space. The insets show an example DNN function (in one-dimensional profile; black line) centred about a sequence of interest with the ground-truth function (dashed line) overlaid. The left inset illustrates the sensitivity of Saliency Maps to non-smooth local function properties. The right inset illustrates the ability of surrogate models to better approximate ground truth.

patterns characteristic of specific biological mechanisms, present within these maps. To provide consistent biological explanations for sequence activity, attribution methods must produce consistent motifs across input sequences that function through shared biological mechanisms.

Established attribution methods for genomic DNNs have major limitations. Different attribution methods use different strategies to quantify position-specific nucleotide effects and can therefore yield different mechanistic explanations^{10–12}. For example, Saliency Maps¹³ quantify nucleotide effects using the gradient of DNN predictions at the sequence of interest, whereas DeepLIFT¹⁴ propagates activation differences between the sequence of interest and a reference sequence. Moreover, the most widely used attribution methods in genomics—including Saliency Maps¹³, DeepLIFT¹⁴, in silico mutagenesis (ISM)¹⁵, SmoothGrad¹⁶, Integrated Gradients¹⁷ and DeepSHAP¹⁸—assume that nucleotide effects on DNN predictions are locally additive. As a result, these attribution methods do not explicitly account for the combinations of genetic interactions (that is, specific epistasis^{19–22}), global nonlinearities (that is, global epistasis (GE)^{23–27}) and heteroscedastic noise²⁶ that are often present in functional genomics data.

In this Article we introduce SQUID (Surrogate Quantitative Interpretability for Deepnets), an interpretability framework for genomic DNNs that overcomes these limitations. SQUID uses surrogate models—simple models with interpretable parameters—to approximate the DNN function within localized regions of sequence space. SQUID applies MAVE-NN²⁶, a quantitative modelling framework developed for analysing multiplex assays of variant effects (MAVEs), to in silico MAVE datasets generated using the DNN as an oracle. SQUID models DNN predictions in a user-specified region of sequence space; this is similar to LIME²⁸, a surrogate modelling approach developed for general DNNs. Unlike LIME, however, SQUID is able to account for important domain-specific outputs of genomic DNNs: global nonlinearities and heteroscedastic noise. Indeed, SQUID can be viewed as a generalization of LIME (see Methods for details).

Benchmarking SQUID against existing attribution methods, we find that SQUID more consistently quantifies the binding motifs of transcription factors (TFs), reduces noise in attribution maps and improves variant-effect predictions. We also find that the domain-specific surrogate models used by SQUID are critical for this improved performance. Finally, we show how SQUID can provide insights into epistatic interactions in *cis*-regulatory elements and can be used to study such interactions both locally and globally in sequence space. SQUID thus provides a new and useful way to interpret genomic DNNs.

Results

Surrogate modelling for genomic DNNs via SQUID

SQUID approximates DNNs in user-specified regions of sequence space using surrogate models that have mechanistically interpretable parameters. The SQUID framework comprises three steps (Fig. 1a): (1) generate an in silico MAVE dataset composed of variant sequences and the corresponding DNN predictions; (2) fit a surrogate model to these in silico data using MAVE-NN²⁶; and (3) visualize and interpret the surrogate model's parameters. This workflow requires the specification of two key analysis parameters: the region of sequence space over which the DNN is to be approximated, and the mathematical form of the surrogate model. By choosing different values for these analysis parameters, users are able to test different mechanistic hypotheses. SQUID assists users by facilitating the use of analysis parameters that have been found in practice to work well in the design and analysis of MAVE experiments.

First, SQUID generates an in silico MAVE dataset. This is done by generating a library of variant sequences, then using the DNN to assign a functional score to each sequence in the library. In this paper, we consider two types of library: local and global. A local library is generated by partially mutagenizing a specific sequence of interest (for example, a genomic *cis*-regulatory sequence). A global library is generated by inserting partially mutagenized versions of a genetic element of interest into random sequences. In what follows, we use a partial mutagenesis rate of 10% per nucleotide (which is common

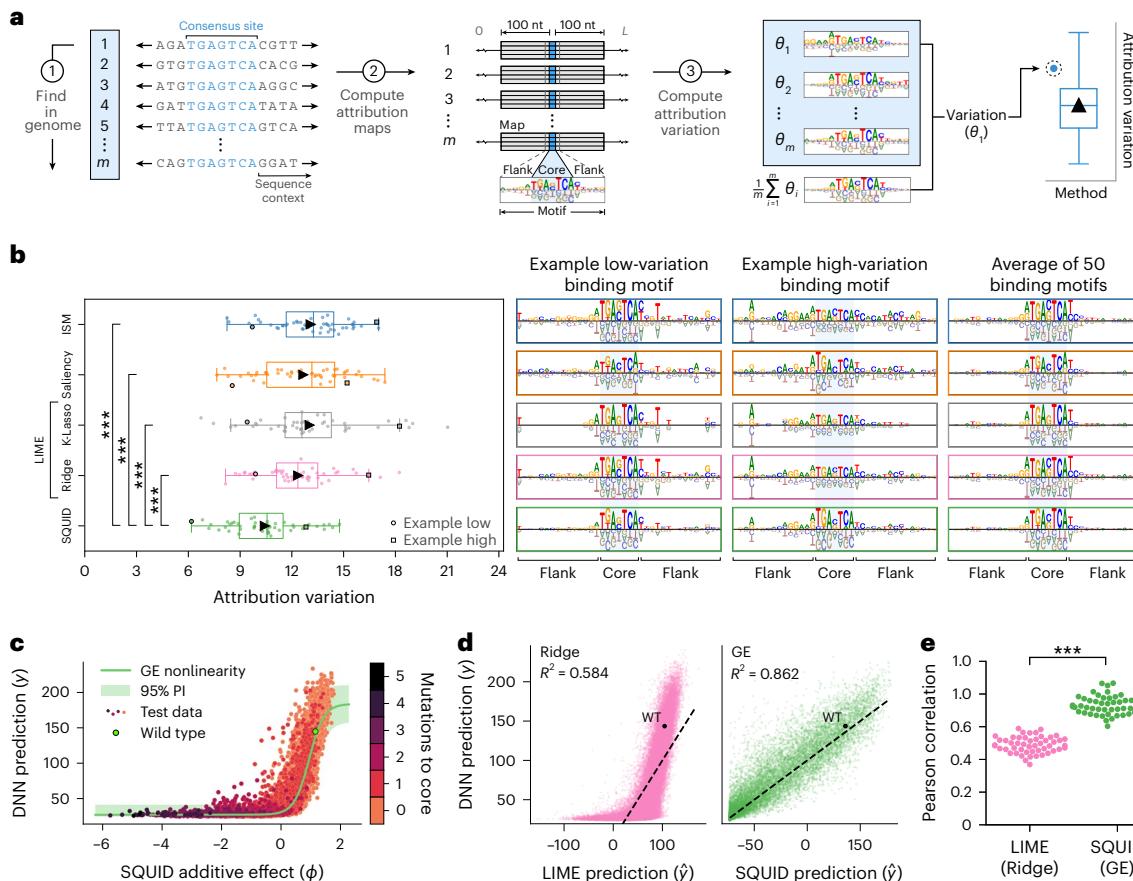


Fig. 2 | Benchmark analysis of attribution methods. **a**, Our benchmark analysis pipeline consisted of three steps: (1) genomic sequences that contained consensus binding sites for a TF of interest were identified; (2) an attribution map spanning the core identified site and 100 nt of flanking sequence on each side was computed for each identified sequence; and (3) a corresponding attribution variation was computed (see Methods for details). **b**, Attribution maps and attribution variation for 7-nt consensus AP-1 binding sites (core sequence TGAGTCA; flank size 15 nt), computed using ResidualBind-32 as the DNN. Left: box plots show attribution variation for $n = 50$ genomic sequences. Triangles represent means; lines represent median, upper quartile and lower quartile; whiskers represent $1.5 \times$ the inter-quartile range. Right: binding motifs observed for two example genomic sequences, together with the ensemble-averaged binding motif. The position of these example binding motifs, representative of low variation and high variation, is marked in the left plot using a circle and

square, respectively. Note: we used $K = 50$ for the K-Lasso implementation of LIME; results for other choices of K are shown in Extended Data Fig. 1.

c, Mutational effects predicted by the DNN versus additive effects predicted by SQUID. Dots correspond to $N = 100,000$ mutagenized sequences, with each dot representing effects of different sets of mutations to a representative genomic sequence. The GE nonlinearity (solid line) and 95% prediction interval (PI; shaded region) inferred by SQUID are shown for comparison. **d**, DNN predictions versus the predictions of two surrogate models for a representative genomic sequence. One surrogate model (GE, right) has a GE nonlinearity; the other (Ridge, left) does not. Dots represent test sequences from the in silico MAVE dataset. Diagonal lines represent equality between DNN and surrogate model predictions. WT, wild-type sequence; R^2 , squared Pearson correlation coefficient. **e**, R^2 values computed as in **d** for 50 different sequences of interest. P -values in **b** and **e** were computed using a one-sided Mann–Whitney U test; *** $P < 0.001$.

in MAVE experiments, for example ref. 29) and libraries comprising 100,000 variant sequences (unless otherwise noted).

Next, SQUID fits a surrogate model to the in silico MAVE dataset. SQUID uses nonlinear surrogate models developed specifically for modelling MAVE data²⁶. These models, called latent phenotype models, have three components (Fig. 1b): a genotype–phenotype (G–P) map, a GE nonlinearity and a noise model. The G–P map projects the input sequence down to a one-dimensional latent phenotype. SQUID, via MAVE-NN, supports the use of additive G–P maps, pairwise-interaction G–P maps and user-defined G–P maps. The GE nonlinearity, which is modelled using a linear combination of sigmoids, maps the latent phenotype to a most-probable DNN prediction. The noise model then describes how actual DNN predictions are expected to scatter around the predicted most-probable value. SQUID supports a variety of noise models, including heteroscedastic noise models based on the skewed t -distribution of Jones and Faddy³⁰ (which we use in this paper). Surrogate model parameters are inferred from in silico MAVE data by maximizing variational

information²⁶ (or, equivalently, log likelihood) using standard stochastic gradient optimization.

The parameters of the G–P map are of primary interest for downstream mechanistic interpretations. The parameters of additive G–P maps quantify single-nucleotide effects, and are readily visualized using the same methods (sequence logos³¹ and heatmaps) normally used for standard attribution maps. The parameters of pairwise-interaction G–P maps quantify epistatic genetic interactions as well as single-nucleotide effects, and can be visualized using block-triangular heatmaps (as in ref. 26).

SQUID improves the quantification of TF binding motifs
A common goal when interpreting attribution maps for a DNA sequence of interest is to identify functional binding sites for TFs. However, the TF binding motifs observed in attribution maps often vary substantially from sequence to sequence. Some of this variation is probably due to underlying biology; for example, TF binding preferences have been experimentally observed to change in response to the binding

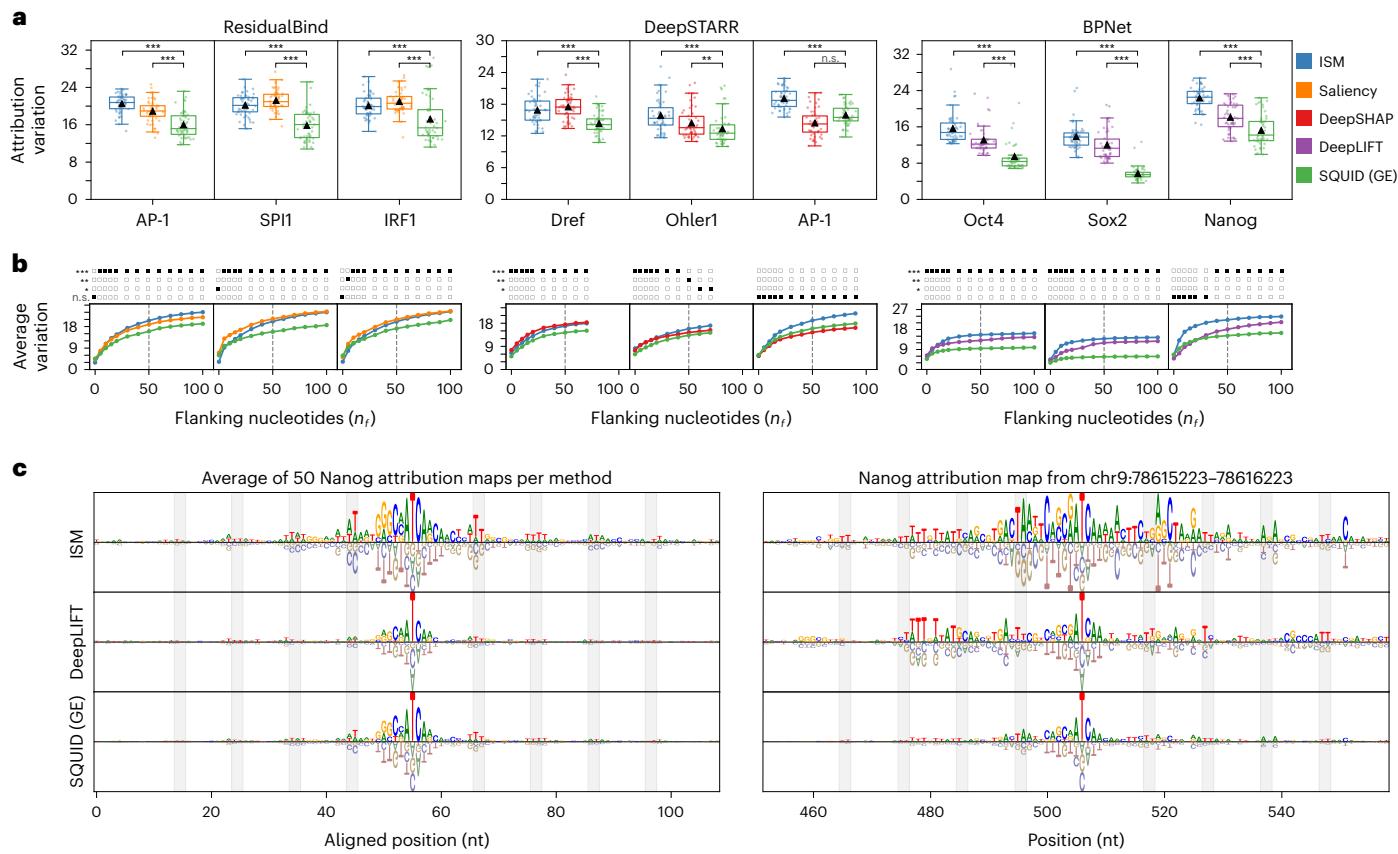


Fig. 3 | Benchmark analysis across TFs, DNNs, and flank sizes. **a**, Attribution variation analysis for various TFs and genomic DNNs. Results are visualized as in Fig. 2b. Each test used $n = 50$ sequences from either the human genome (for ResidualBind-32), the mouse genome (for BPNet) or the fly genome (for DeepSTARR), together with consensus TF binding site listed in Supplementary Table 1, and flanked by 50 nt of DNA. **b**, Mean attribution scores computed as in **a**, but using variable lengths of flanking DNA. The boxes above each plot indicate

the largest (that is, least significant) P value, computed as in Fig. 2b. **n.s.**, $P \geq 0.05$; * $0.01 \leq P < 0.05$; ** $0.001 \leq P < 0.01$; *** $P < 0.001$. Dashed line, flank size used in **a**. **c**, Binding motifs computed for Nanog via attribution analysis of BPNet. Left: attribution maps averaged across 50 mouse loci containing (and centred on) the consensus Nanog binding site AGCCATCAA. Right: attribution map observed for a single such locus. Grey bars indicate 10.5 nt periodicity on either side of the consensus Nanog binding site.

of other TFs nearby^{32–34}. However, variation in TF binding motifs can be exacerbated by the specific ways in which attribution methods quantify the behaviour of the DNN function in localized regions of sequence space (Fig. 1c). The consistency of binding motifs observed in attribution maps across different genomic sequences therefore provides a way to quantify and compare the performance of different attribution methods.

To benchmark the consistency of binding motifs identified by different attribution methods, we identified sequences in the human genome that contain putative TF binding sites. Specifically, for each TF, we identified genomic instances of the consensus TF binding site sequence at which a strong motif was observed according to a baseline attribution method. We then aligned these genomic sequences about their putative binding sites and computed, for each genomic sequence, an attribution map that spans the core putative binding site as well as a specified amount of flanking DNA on either side. Each attribution map was then normalized to control for locus-to-locus variation in overall motif scale (scale was substantial). Finally, we calculated the Euclidean distance between the vector of normalized attribution scores for individual sequences and the mean vector of normalized attribution scores. We refer to this distance as the ‘attribution variation’ (Fig. 2a; see Methods for details).

We first applied this benchmarking pipeline to the human TF AP-1 using ResidualBind-32 (ref. 35), a genomic DNN that predicts chromatin accessibility in human cell lines. We compared SQUID with two commonly used attribution methods that had previously been used in ref. 35

to analyse ResidualBind-32: ISM and Saliency Maps. We also compared SQUID with two different implementations of LIME. One implementation, Ridge, is similar to SQUID but lacks a GE nonlinearity and uses a homoscedastic Gaussian noise model. The other implementation, K-Lasso, uses the specific algorithm described in ref. 28 and requires users to specify the number of non-zero features in the attribution map.

We found that, across different genomic sequences that contain the consensus AP-1 binding site TGAGTC, SQUID recovers AP-1 binding motifs that have markedly lower attribution variation than the binding motifs recovered by ISM, Saliency Maps and both implementations of LIME (Fig. 2b, left, and Extended Data Fig. 1). This result is supported by the examination of attribution maps of individual sequences (Fig. 2b, right). Compared with the attribution maps provided by other methods, the attribution maps provided by SQUID exhibit core regions with greater similarity to the ensemble-averaged motif, and flanking regions with reduced (probably non-biological) scores. Visual examination of individual PA-1 motifs suggest that some of the observed motif variation was biological in nature and that in such cases SQUID was better able than other methods to discern functional variation from noise in these cases. Moreover, the motifs provided by SQUID were largely robust to the choice of SQUID analysis parameters (Extended Data Fig. 2). We conclude that SQUID quantifies the AP-1 binding motif from ResidualBind-32 more consistently than does ISM, Saliency Maps or LIME.

We next investigated whether the surrogate models for AP-1 identified by SQUID benefitted from incorporating a GE nonlinearity.

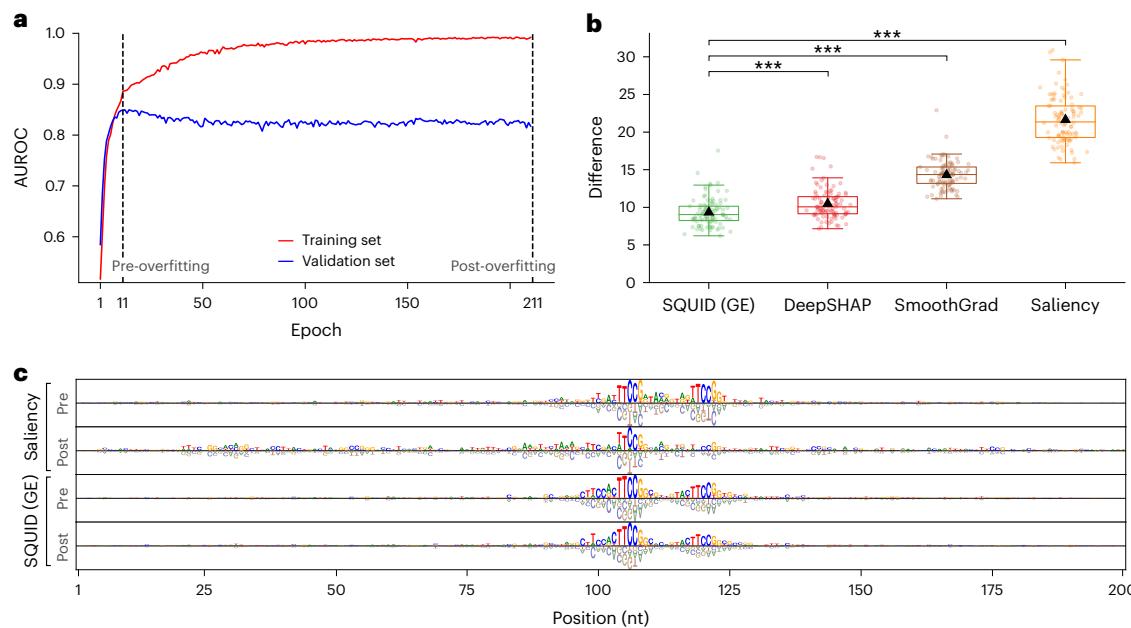


Fig. 4 | Attribution method performance during benign overfitting.

a, DNN performance as a function of training epoch. DNN was a three-layer convolutional neural network trained to classify 200-nt ChIP-seq peaks for the human TF GABPA. Tests used DNN parameters from epoch 11 (pre-overfitting) and epoch 211 (post-overfitting). **b**, Differences between attribution maps for the DNN with pre-overfitting parameters versus post-overfitting parameters, as

quantified by the Euclidean distance between attribution map vectors.

Box plot results are shown for $n = 100$ genomic sequences in the ChIP-seq peak test set. Data are visualized and statistical tests were performed as in Fig. 2b.

c, Attribution maps obtained for a representative test sequence using pre-overfitting and post-overfitting DNN parameters.

We found that SQUID quantifies the AP-1 binding motif better when using GE regression than when using the Ridge implementation of LIME (Fig. 2b). Plotting the effects that mutations in a representative genomic sequence have on DNN predictions, we found that virtually all combinations of two or more mutations to the core 7-nt AP-1 site reduced DNN predictions to near-background levels (Fig. 2c). Moreover, the GE nonlinearity learned by SQUID as part of the surrogate model accurately recapitulated this saturation effect. By contrast, surrogate modelling using the Ridge implementation of LIME failed to capture this saturation effect (Fig. 2d,e). This finding demonstrates that surrogate modelling of genomic DNNs can benefit from using our domain-specific surrogate models, as opposed to the linear surrogate models that are standard in other fields (for example, computer vision²⁸).

We then expanded our analysis to other TFs and to other genomic DNNs (DeepSTARR³⁴ and BPNet⁵; Supplementary Table 1). In each benchmark analysis, we compared SQUID with the attribution methods (ISM, Saliency Maps, DeepLIFT or DeepSHAP) used in the original study describing the DNN being modelled. We found that the attribution maps provided by SQUID consistently yielded binding motifs with markedly lower attribution variation than the binding motifs provided by the other attribution methods (Fig. 3a). These results were robust to the amount of flanking DNA used when computing attribution variation (Fig. 3b). We also found that strong GE nonlinearities were pervasive in the surrogate models inferred by SQUID for the genomic sequences tested (Extended Data Fig. 3). These results suggest that SQUID, quite generally, quantifies TF binding motifs more consistently than do competing attribution methods. These findings also suggest that modelling GE nonlinearities and heteroscedastic noise is important for the accurate surrogate modelling of genomic DNNs.

The ability of SQUID to better quantify TF binding motifs is exemplified in Fig. 3c. Shown are attribution maps for the mouse TF Nanog, computed using BPNet, a consensus binding site of AGCCATCAA, and 50 nt of flanking DNA. When attribution maps are averaged across genomic loci, ISM, DeepLIFT and SQUID all produce attribution maps that reveal both the Nanog binding motif and a more subtle preference

for periodically spaced AT-rich sequences in the flanking DNA (which probably reflects the interaction of nucleosomes with the DNA double helix). However, the attribution maps for specific genomic sequences often exhibit many other noticeable features. Some of these features are likely to be biological, as TFs often bind in clusters³⁰. However, many features appear to be spurious and probably reflect non-biological noise in the attribution map. The attribution map provided by SQUID appears to exhibit less noise in the sequences that directly flank core motifs than the attribution maps given by ISM or DeepLIFT. Similar observations hold in analyses of other TFs (Extended Data Fig. 4). These findings raise the possibility that the increased consistency of binding motifs identified by SQUID is due, at least in part, to the reduction of non-biological noise in attribution maps.

SQUID reduces noise in attribution maps

Non-biological noise in attribution maps can arise from the inherent roughness in the DNN function, a phenomenon that is often associated with benign overfitting^{37,38}. Benign overfitting is of particular concern for Saliency Maps, since it can adversely affect the accuracy of DNN gradients^{39–41}. Benign overfitting is similarly expected to impact other attribution methods, since these methods essentially quantify how small changes in sequence space affect DNN predictions^{10,40}. We reasoned that, because SQUID integrates information over an extended region of sequence space, the attribution maps provided by SQUID might be less noisy than the maps provided by other attribution methods.

We therefore asked whether SQUID can reduce the attribution map variation caused by benign overfitting. To answer this question, we trained a DNN to classify chromatin immunoprecipitation followed by sequencing (ChIP-seq) peaks for the human TF GABPA (Methods) and saved DNN parameters both before benign overfitting (that is, using early stopping) and after benign overfitting (Fig. 4a). Benign overfitting is apparent from the plateau in validation-set performance that occurs when near-perfect classification performance is achieved on the training set⁴². We then selected 100 random sequences from the test set, and for each sequence and each of four attribution methods (SQUID,

Table 1 | Attribution method performance on a zero-shot variant-effect prediction task

DNN		Average Pearson correlation				Statistical significance		
Architecture	Task	Saliency	ISM	SQUID (Ridge)	SQUID (GE)	GE versus Saliency	GE versus ISM	GE versus Ridge
Enformer	DNASE	0.2977	0.4498	0.4486	0.4801	***	**	**
Basenji-32	ATAC-seq	0.2727	0.3575	0.3700	0.4036	***	***	**
ResidualBind-32	ATAC-seq	0.2846	0.3388	0.3567	0.3912	**	***	*

Performance of four attribution methods, applied to three DNNs, on the CAG15 variant-effect prediction challenge⁴⁷. Numbers indicate the average correlation observed across the 15 genomic loci assayed by MPRA in ref. 47. Bold text indicates the best-performing method for each DNN. P values report results from a paired one-sided Wilcoxon signed-rank test on the 15 locus-specific correlation values: * $0.01 \leq P < 0.05$; ** $0.001 \leq P < 0.01$; *** $P < 0.001$.

DeepSHAP, SmoothGrad and Saliency Maps), quantified the differences between the attribution maps obtained using the pre-overfitting DNN parameters versus the post-overfitting DNN parameters. Figure 4b shows that SQUID provided attribution maps that changed substantially less over the course of benign overfitting. Figure 4c illustrates this behaviour for a representative sequence. These results support the hypothesis that the attribution maps provided by SQUID are more robust to the adverse effects of benign overfitting than are other attribution methods. These results also suggest that SQUID, more generally, yields attribution maps that have lower noise than maps computed by other attribution methods.

SQUID better identifies putative weak TF binding sites

Weak TF binding sites play critical roles in eukaryotic gene regulation^{43–45}. Functional signals from weak binding sites, however, can be difficult to distinguish from noise in attribution maps. Having shown that SQUID reduces noise in attribution maps relative to other attribution methods, we hypothesized that SQUID would also better identify weak yet functional TF binding sites. To test this hypothesis, we quantified how well attribution maps generated for putative weak TF binding sites matched the TF binding motifs identified in Fig. 3a. For each TF of interest, we randomly selected 150 putative binding sites in the genome having zero, one or two mutations relative to the consensus binding site (Supplementary Table 1). We also recorded the genomic sequence containing each selected site padded by 50 nt of flanking DNA on either side. We then computed the score assigned to each selected site using a TF-specific position weight matrix (PWM)⁴⁶, as in Extended Data Fig. 5a. Different attribution methods were used to compute attribution maps for each genomic sequence, after which each sequence was assigned an attribution variation value quantified by the Euclidean distance between the attribution map and the corresponding ensemble-averaged attribution map from Fig. 3a. Extended Data Fig. 5b shows that, as expected, the resulting variation for all attribution methods increased as PWM score decreased. However, the attribution variation observed for SQUID was consistently lower than the attribution variation observed for competing methods. This finding is confirmed by visually examining attribution maps for selected sites (Extended Data Fig. 6). We conclude that SQUID is better than competing methods at identifying signatures of TF binding at weak yet functional TF binding sites.

SQUID improves zero-shot prediction of SNV effects

A major goal of genomic DNNs and their attribution methods is to predict which genetic variants are pathogenic. Having observed that SQUID reduces noise in attribution maps, mitigates the adverse effects of benign overfitting and better identifies TF binding motifs, we hypothesized that SQUID would provide improved variant-effect predictions. To test this hypothesis, we used SQUID and other attribution methods to predict the effects that single-nucleotide variants (SNVs) have on the activity of *cis*-regulatory elements for 15 different disease-associated loci in the human genome, loci for which the effects of SNVs had been measured using massively parallel reporter assays (MPRAs)^{47,48}.

In technical terms, this is a zero-shot prediction task, since none of the methods tested were trained on MPRA data. At each locus, the Pearson correlation coefficient was computed between the attribution scores and measured SNV effects. We performed this benchmark analysis using three genomic DNNs previously reported for predicting regions of open chromatin: Enformer², Basenji-32 (ref. 49) and ResidualBind-32 (ref. 35). We found that SQUID yielded substantially higher correlations between predicted and measured SNV effects across the 15 assayed loci (Table 1 and Supplementary Fig. 1). The results suggest that attribution maps provided by SQUID are generally able to predict variant effects better than the attribution maps provided by competing methods. The improved performance of SQUID relative to ISM, which directly uses DNN predictions to quantify SNV effect sizes, suggests that surrogate modelling may provide a general way of improving the variant-effect predictions of genomic DNNs themselves.

SQUID illuminates epistatic interactions

Understanding the role of epistatic interactions within gene regulatory sequences is a major goal in the study of *cis*-regulatory codes. An important advantage of surrogate modelling over other DNN interpretability approaches is that surrogate models with different mathematical forms can be used to answer different questions about DNN behaviour. To test whether SQUID could be used to study epistatic interactions in *cis*-regulatory sequences, we implemented a surrogate model that describes all possible pairwise interactions between nucleotides within a sequence (in addition to the additive contributions from individual nucleotides). We then used this pairwise-interaction model to quantify the effects of pairs of putative AP-1 binding sites learned by ResidualBind-32³⁵. We identified 50 genomic sequences having pairs of putative AP-1 binding sites spaced 4–20 nt apart, used SQUID to infer the parameters of the pairwise-interaction model about each of the 50 sequences, and then aligned and averaged the values of the pairwise-interaction parameters corresponding to both inter-site and intra-site interactions across the 50 sequences (Fig. 5a).

The results are shown in Fig. 5b,c. We found that pairwise-interaction models consistently performed better on test data than additive surrogate models inferred in a similar manner (Fig. 5b). This was true in both the presence and absence of a GE nonlinearity (both additive and pairwise-interaction models benefitted from having a GE nonlinearity). The pairwise-interaction models thus provide more accurate approximations of the DNN, suggesting that the pairwise-interaction parameters in these models are likely to be meaningful. Examining the resulting context-averaged pairwise-interaction parameters (Fig. 5c), we observed strong positive intra-site interactions and strong negative inter-site interactions for critical mutations, that is, mutations away from the preferred nucleotide at any of the six most selective positions of the AP-1 motif. Intuitively, the positive interactions between a pair of critical mutations within the same AP-1 site arise because a single critical mutation is sufficient to abrogate AP-1 binding. Thus, a second critical mutation within the same site has no effect, rather than the additional negative effect that would be predicted by an additive model. By contrast, the negative inter-site interactions

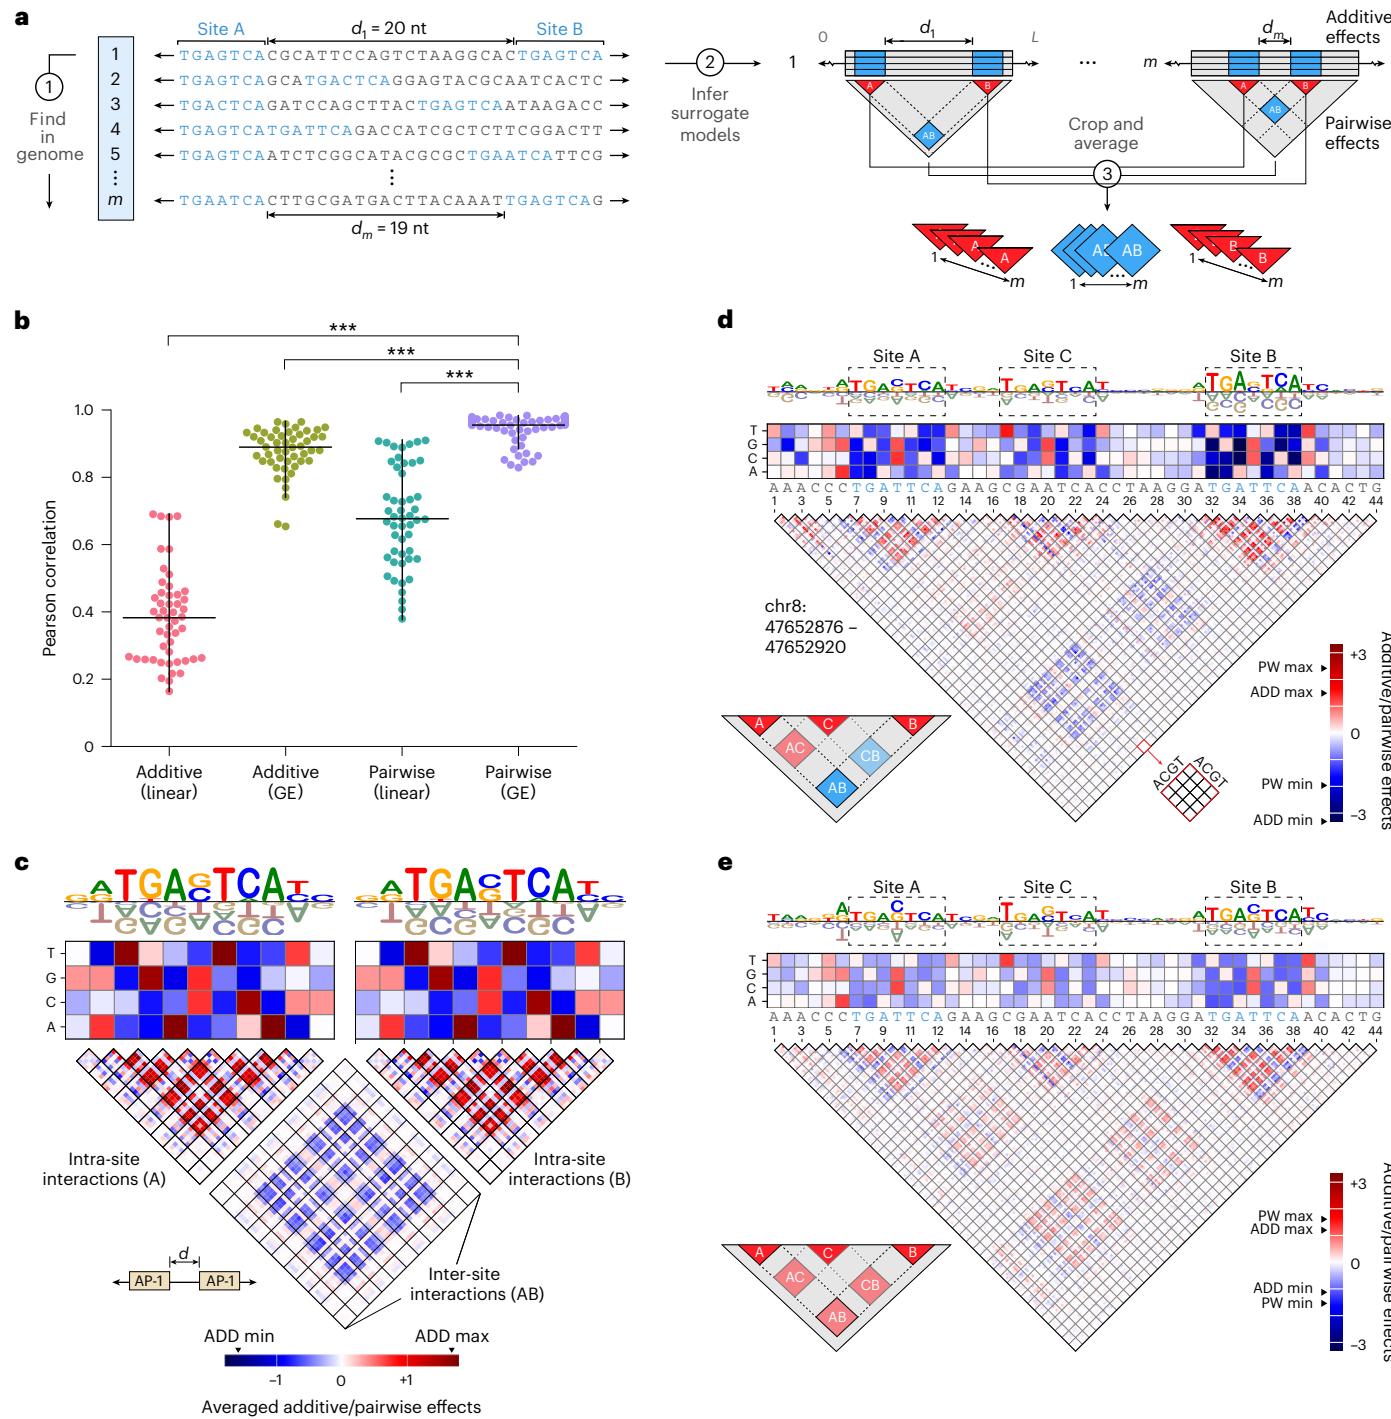


Fig. 5 | SQuID captures epistatic interactions. **a**, Pipeline for analysing context-averaged epistatic interactions. The pipeline consists of three steps: (1) pairs of consensus binding sites (A and B) are identified in genomic sequences; (2) pairwise-interaction models are inferred for each identified sequence; and (3) the values of surrogate model parameters describing intra-site interactions (A, B) and inter-site interactions (AB) are aligned and averaged across sequence contexts. **b**, Performance of surrogate models for $n = 50$ genomic sequences having two consensus AP-1 binding sites each. Swarm plot results are shown for both additive and pairwise-interaction models, each with (GE) and without (linear)

a GE nonlinearity. Correlation values were computed between surrogate model predictions and DNN predictions on *in silico* test data. Overlaid lines represent the median, upper and lower quartiles. P values were computed using the one-sided Mann–Whitney U test. *** $P < 0.001$. **c**, Surrogate model parameters quantifying intra-site and inter-site interactions, aligned and averaged across the 50 genomic sequence contexts. Pairwise-interaction models having GE nonlinearities were used in this analysis. **d,e**, Parameters of pairwise-interaction models, either with (**d**) or without (**e**) a GE nonlinearity, determined for a genomic locus having three putative AP-1 binding sites. ADD, additive model; PW, pairwise model.

between critical mutations indicate a degree of redundancy between nearby AP-1 binding sites, since mutations that abrogate both AP-1 sites result in lower activity than is expected from adding the effects of abrogating each AP-1 site individually⁵⁰.

While these averages are informative about overall trends, analyses of specific sequences can provide additional insight. In the analysis above, we observed one sequence that contained three putative AP-1 binding sites, with different combinations of these sites displaying

either positive or negative epistatic interactions (Fig. 5d). This example shows that context-specific interactions can be complex and that other factors like binding site orientation and spacing may play an important role⁵¹. We note that including the GE nonlinearity in the pairwise-interaction model was essential for identifying these complex interactions. When a linear pairwise-interaction model was used, we instead observed positive interactions of similar magnitude between every pair of sites (Fig. 5e). The reason is that, in the linear pairwise model, the pairwise-interaction parameters are co-opted to describe a global nonlinearity instead of the nucleotide-specific interactions they are intended to model. Specifically, for this genomic sequence, the DNN exhibits a rapidly saturating effect of removing AP-1 binding sites on the functional score (Extended Data Fig. 7), such that disrupting a second binding site has a much smaller effect on DNN predictions than would be predicted on the basis of the effect of disrupting just one of the three sites. We thus see that modelling GE nonlinearities is essential when interpreting genomic DNNs in terms of epistatic interactions between genetic elements.

SQUID supports global DNN interpretations

In the above analyses, SQUID used *in silico* MAVE libraries generated by partially mutagenizing a specific sequence of interest in the genome. As a result, the surrogate models inferred by SQUID provided DNN interpretations that are limited to localized regions of sequence space. In previous work, however, we proposed a DNN interpretation method called global importance analysis⁵², in which the DNN is evaluated on completely random sequences containing embedded genetic elements, such as putative TF binding sites. We therefore investigated whether SQUID could provide useful global DNN interpretations using a modified version of the sequence libraries used in global importance analysis.

First we asked whether SQUID could provide global interpretations for individual TFs. For each TF of interest, we generated an *in silico* MAVE library by embedding partially mutagenized versions of the consensus TF binding site within random DNA (Extended Data Fig. 8a). We then used this *in silico* MAVE library to infer an additive surrogate model. This approach was applied to the four mouse TFs (Oct4, Sox2, Klf4 and Nanog) modelled by BPNet⁵. Extended Data Fig. 8b shows that the resulting additive surrogate models revealed prominent motifs having low background, which matched the known sequence preferences of the four TFs. In the case of Nanog, this analysis also revealed periodic secondary motif features in flanking DNA nearly identical to those in Fig. 3c. We conclude that global surrogate modelling with SQUID can provide accurate characterizations of TF binding motifs independent of specific genomic sequence contexts.

We next investigated whether SQUID could provide global insights into the epistatic interactions between pairs of TF binding sites. We created an *in silico* MAVE library in which partially mutagenized version of the Nanog and Sox2 consensus binding sites were inserted into random DNA sequences a fixed distance apart (Extended Data Fig. 8c). We then used SQUID to infer a pairwise-interaction model based on this *in silico* MAVE library. Extended Data Fig. 8d shows that, similar to the findings in Fig. 5, critical mutations within Nanog and Sox2 binding sites exhibited positive intra-site epistatic interactions and negative inter-site epistatic interactions. We then repeated this analysis for Nanog and Sox2 binding sites separated 0 to 32 nt and observed that the strength of inter-site epistatic interactions varied in a sinusoidal manner consistent with the periodicity of the DNA double helix (Extended Data Fig. 8e). We conclude that global surrogate modelling with SQUID can provide useful characterizations of epistatic interactions between TFs in a manner that is independent of genomic sequence context.

Discussion

Here we introduced SQUID, a framework for interpreting genomic DNNs. SQUID uses surrogate models to approximate DNN functions in user-defined regions of sequence space. The parameters of these

surrogate models can then be mechanistically interpreted: additive surrogate models can be interpreted as attribution maps, and pairwise-interaction surrogate models can be interpreted as quantifying epistatic interactions. Applying SQUID to a variety of genomic DNNs, we observed that the attribution maps obtained by SQUID more robustly identify TF binding motifs and provide better variant-effect predictions than the attribution maps obtained using other DNN interpretability methods. We also observed that SQUID is able to quantify epistatic interactions that are otherwise obscured by global nonlinearities in the DNN.

SQUID works by using the DNN of interest as a forward simulator of MAVE experiments. SQUID then uses the MAVE-NN modelling framework²⁶ to infer surrogate models from the resulting *in silico* data. This surrogate modelling approach has multiple important advantages over standard DNN interpretability methods.

First, SQUID is model agnostic: it does not require access to the parameters or gradients of the DNN. Rather, SQUID simply uses the DNN of interest as a black-box oracle. This allows SQUID to be applied to arbitrary genomic DNNs regardless of their computational implementation.

Second, SQUID smooths over fluctuations in DNN predictions that are likely to reflect noise rather than biological function. This is because the parameters of the surrogate models that SQUID infers are fit to DNN predictions over an extended region of sequence space. Our results showed that this increased smoothness can cause the surrogate models to have improved accuracy relative to the parent DNN.

Third, SQUID leverages domain-specific knowledge to improve the utility of surrogate models. While the use of surrogate models for interpreting DNNs has been established previously in other fields, these applications typically use linear models²⁸. By using the latent phenotype models supported by MAVE-NN, the surrogate models inferred by SQUID explicitly account for global nonlinearities and heteroscedastic noise, both of which are ubiquitous in functional genomics data and in genomic DNNs. By explicitly accounting for these influences, SQUID is able to remove their confounding effects on the inferred surrogate model parameters.

Finally, SQUID supports a variety of surrogate models. In this work, we demonstrated the use of additive and pairwise-interaction models, but SQUID also supports use of surrogate models having user-specified mathematical form. SQUID thus has the ability to infer surrogate models that reflect higher-order epistatic interactions^{21,22,25,53}, or specific biophysical hypotheses (for example, thermodynamic models^{54–58} or kinetic models^{59–62}).

Although nonlinearities and genetic interactions can be identified by observing whether changes to sequence cause changes to additive attribution maps (for example, maps provided by Saliency Maps, DeepLIFT or DeepSHAP), explicitly accounting for these nonlinearities and genetic interactions within the surrogate model allows SQUID to quantify these effects in a natural way. One limitation to modelling epistatic interactions in SQUID using pairwise-interaction surrogate models is that the number of such interactions, and thus the number of parameters that must be inferred, increases quadratically with sequence length. We found, in practice, that the inference of pairwise-interaction models presented computational difficulties when input sequences were longer than approximately 100 nt. The number of parameters in additive models increases only linearly with sequence length, and such models are often more practical to implement on long sequences as a result.

The biggest drawback of SQUID relative to other DNN interpretability methods (other than LIME, which has the same drawback) is its higher computational demands. This is due to the need for many forward passes through the DNN during simulation of the *in silico* MAVE dataset, as well as the need to fit the surrogate model parameters to these simulated data (see Extended Data Table 1 for typical SQUID computation times). We therefore suggest that SQUID is likely to be more useful for in-depth analyses of specific sequences of interest.

(for example, disease-associated loci) rather than large-scale genome-wide analyses. In this context, however, the advantages of SQUID described above—especially SQUID’s ability to support surrogate models of arbitrary mathematical form—enables new ways of biologically interpreting genomic DNNs.

Methods

The SQUID framework

An overview of SQUID’s workflow is given in Supplementary Fig. 2. Briefly, SQUID takes as input a sequence of interest and a specified surrogate model. An in silico MAVE dataset is generated with the InSilicoMAVE object, with specifications of the mutagenesis strategy given by a Mutagenizer object, and processed model predictions given by a Predictor object. The in silico MAVE dataset is then fit with surrogate models that are defined as objects in squid/surrogate_zoo.py.

- **Mutagenizer.** An in silico MAVE dataset is generated by sampling a library of sequences using random partial mutagenesis of a sequence of interest. We modulate the size of the sequence-space region from which this library is drawn using two hyperparameters: the sequence that defines the region of interest, which has length L , and the mutation rate r . The resulting number of mutations in each individual sequence is a Poisson distributed random variable having mean Lr . squid/mutagenizer.py contains objects that apply the chosen mutagenesis strategy, with the object RandomMutagenesis executing the random partial mutagenesis in this study.
- **Predictor.** SQUID currently requires that DNNs provide scalar-valued outputs. However, some genomic DNNs output high-dimensional predicted profiles, not scalar predictions. For example, ResidualBind-32 predicts 15 chromatin accessibility profiles of length 64, each profile corresponding to a different cell type and each profile element representing a binned position with a resolution of 32 nt. squid/predictor.py contains objects, including ScalarPredictor and ProfilePredictor, that reduce model predictions to scalar values. For profile-based predictions, SQUID also offers an approach to reduce profiles using principal component analysis (PCA). Specifically, profiles are projected onto their first principal component, with sign chosen so that the wild-type sequence has higher-than-average score (Supplementary Fig. 3).
- **InSilicoMAVE.** The object InSilicoMAVE, defined in squid/mave.py, is a data structure that takes the mutagenizer and predictor objects and generates the in silico MAVE dataset for a sequence of interest.
- **Surrogate Zoo.** The squid/surrogate_zoo.py module currently offers a linear model (SurrogateLinear) as well as models based on MAVE-NN (SurrogateMAVENN). In SurrogateLIME, the ‘K-Lasso’ approach is implemented²⁸. First, sklearn regression is performed using linear_model.LassoLars to model the entire regularization path over all values of α (ref. 63). We then select the model associated with the appropriate α that gives the desired number of non-zero features, K , that is, weights. These K features are then fit using least squares linear regression. In SurrogateRidgeCV, a linear model is computed using sklearn ridge regression with iterative fitting along a regularization path, where the best model is selected by cross-validation (linear_model.RidgeCV)⁶³. SurrogateMAVENN supports models based on the built-in modelling capabilities of MAVE-NN. The mathematical forms of these models are described in detail in ref. 26

Surrogate models

The surrogate models supported by MAVE-NN for use in SQUID comprise three parts (Fig. 1b): a G-P map, a GE nonlinearity and a noise model. Here, we summarize the mathematical forms of these model

components. In what follows, x represents a sequence of interest, x_{lc} is a one-hot encoding of x (that is, is equal to 1 if x has character c at position l and is equal to 0 otherwise), and y is the DNN-predicted scalar activity of x .

- The G-P map, $f(x, \theta)$, which has parameters θ , maps a sequence x to a latent phenotype ϕ . There are two types of G-P map: additive G-P maps and pairwise-interaction G-P maps. Additive G-P maps have a constant parameter θ_0 , additive parameters θ_{lc} , and the mathematical form

$$f_{\text{additive}}(x, \theta) = \theta_0 + \sum_{l=1}^L \sum_c \theta_{l:c} x_{l:c}.$$

Pairwise-interaction G-P maps have a constant parameter, additive parameters, pairwise-interaction parameters $\theta_{l:c,l':c'}$ and the mathematical form

$$f_{\text{pairwise}}(x, \theta) = \theta_0 + \sum_{l=1}^L \sum_c \theta_{l:c} x_{l:c} + \sum_{l=1}^{L-1} \sum_{l'=l+1}^L \sum_{c,c'} \theta_{l:c,l':c'} x_{l:c} x_{l':c'}.$$

- The GE nonlinearity, $g(\phi)$, maps the latent phenotype ϕ to a most-probable scalar DNN prediction \hat{y} . By default, $g(\phi)$ is defined to be an over-parameterized linear combination of hyperbolic tangent sigmoids. In this paper we use 50 sigmoids, corresponding to 50 hidden nodes in the neural network formulation of $g(\phi)$. This is the default number of sigmoids used by MAVE-NN, and was chosen on the basis of previous analyses of diverse MAVE datasets²⁶. We found that TF motifs inferred by SQUID, as well as corresponding R^2 values, were highly robust to the number of hidden nodes (Extended Data Table 2). For SQUID analyses performed without a GE nonlinearity, $g(\phi)$ is defined to be a linear function of ϕ .
- The noise model, $p(y|\hat{y})$, describes the expected distribution of DNN predictions y about the most-probable prediction \hat{y} . The noise model can be defined using a Gaussian distribution, a Student’s t -distribution, or the skewed t -distributed of Jones and Faddy³⁰. The shape parameters of this distribution can either be independent of \hat{y} (for homoscedastic noise) or a polynomial function of \hat{y} (for heteroscedastic noise).

LIME

LIME is a surrogate modelling approach for interpreting general DNNs. Here, we describe the general LIME framework, how SQUID generalizes LIME and how the two implementations of LIME used for the benchmarking results in Fig. 2 and Extended Data Fig. 1 work.

Let x denote a one-hot encoded sequence input, $y(x)$ denote the prediction of the DNN of interest, $\pi(x)$ denote a probability distribution on sequence space, $f(x, \theta)$ denote the surrogate model used to approximate y , and θ denote the parameters of the surrogate model. We note that ref. 28 distinguishes between the input x fed to the DNN and the ‘interpretable input representation’ x' used as input to the surrogate model inferred by LIME. In our genomics applications, however, we take both x and x' to be the standard one-hot encoding of sequence. We therefore omit this distinction in what follows.

LIME determines values for surrogate model parameters via

$$\theta^* = \underset{\theta}{\operatorname{argmin}} L_{\text{LIME}}(y, \theta, \pi) + \Omega_{\text{LIME}}(\theta),$$

where L_{LIME} is a loss function and Ω_{LIME} is a regularization penalty on θ . Following ref. 28, we adopt the weighted L_2 loss function

$$L_{\text{LIME}} = \sum_x \pi(x) [y(x) - f(x, \theta)]^2,$$

which is estimated by randomly sampling inputs x according to the probability distribution $\pi(x)$.

SQUID can be thought of as a generalization of LIME in which the loss function has latent parameters that must be inferred simultaneously with the parameters of the surrogate model. As in LIME, the G–P map $f(x, \theta)$ serves as the surrogate model. But unlike LIME, the SQUID loss function is determined by a measurement process $p_{\text{MP}}(y|f, \eta)$, which represents a probability distribution over DNN predictions y conditioned on the prediction f of the G–P map, and which has its own set of latent parameters (denoted by η). SQUID simultaneously determines values for θ and η via

$$\theta^*, \eta^* = \underset{\theta, \eta}{\operatorname{argmin}} \ L_{\text{SQUID}}(y, \theta, \eta, \pi) + \Omega_{\text{SQUID}}(\theta, \eta),$$

where Ω_{SQUID} is a regularization penalty on both θ and η , and the loss function L_{SQUID} is given by

$$L_{\text{SQUID}} = \sum_x \pi(x) \log [p_{\text{MP}}(y(x)|f(x, \theta), \eta)].$$

In terms of the GE nonlinearity $\hat{y}(f, \eta)$ and the noise model $p_{\text{noise}}(y|\hat{y}(f, \eta))$, the measurement process is given by

$$p_{\text{MP}}(y|f, \eta) = p_{\text{noise}}(y|\hat{y}(f, \eta), \eta). \quad (1)$$

In ‘Main’, we consider two different implementations of LIME: Ridge and K-Lasso. The Ridge implementation is equivalent to SQUID when the GE nonlinearity is taken to be the identity function and the noise model is taken to be a Gaussian with fixed variance. In this implementation, we also take Ω_{LIME} to be an L_2 penalty with scaling coefficient determined by cross-validation. The K-Lasso implementation is equivalent to the specific algorithm described in ref. 28, if $\pi(x)$ is taken to be the distribution in sequence space resulting from a Poisson mutation process. In this implementation, the user selects a value K for the number of non-zero one-hot features (as DNA sequence of length L having $4L$ such features). Lasso regression is then used to select the K specific features for which θ is allowed to be non-zero, and the specific values of these non-zero elements of θ are inferred by least squares.

Deep learning models

This study used six DNNs: ResidualBind-32 (ref. 35), Basenji-32 (ref. 35), DeepSTARR³⁴, Enformer², BPNet⁵ and a baseline convolutional neural network (CNN) that predicts ChIP-seq data for the human TF GABPA. Here, we briefly describe each DNN and how that DNN was used in our study to compute a prediction y for each sequence x when generating *in silico* MAVE data.

- ResidualBind-32 predicts assay for transposase-accessible chromatin with sequencing (ATAC-seq) profiles across 15 human cell lines³⁵. ResidualBind-32 takes as input a DNA sequence of length 2,048 nt and outputs 15 profiles (one for each cell line) where each profile comprises 64 bins, with each bin spanning 32 nt. The published ResidualBind-32 parameters were used to compute these profiles. In our attribution variation analyses, y was set equal to the sum of predicted binned ATAC-seq signals over all 64 bins for the single output channel corresponding to the cell type most associated with the TF of interest (Supplementary Table 1). In our variant-effect analysis, y was set equal to the sum of predicted binned ATAC-seq signals over all 64 bins, using a profile averaged across all 15 output channels.
- Basenji-32 predicts ATAC-seq profiles across 15 human cell lines⁴⁹. The input and output of Basenji-32 is identical to that of ResidualBind-32, and the published Basenji-32 parameters were used to compute predicted ATAC-seq profiles. Analyses performed using Basenji-32 were performed the same way as for ResidualBind-32.
- DeepSTARR predicts *Drosophila* enhancer activity as assayed by UMI-STARR-seq³⁴. DeepSTARR takes as input a DNA

sequence of length 249 nt and outputs two scalar-valued predictions for enhancer activity for developmental (Dev) and housekeeping (Hk) regulatory programs. The published DeepSTARR parameters were used to predict enhancer activity. In each analysis, y was computed using the regulatory program most associated with the TF of interest (Supplementary Table 1).

- Enformer predicts many different types of functional genomic track (for example, ChIP-seq, DNase-seq, ATAC-seq and cap analysis of gene expression) across the human and mouse genomes². Enformer takes as input a DNA sequence of length 393,216 nt and (for humans) outputs 5,313 profiles (one for each track) where each profile comprises 128 bins, each bin spanning 32 nt, representing the central 114,688 nt of the input sequence. The published Enformer parameters were used to compute these profiles. In our variant-effect analysis, we used human predictions, where y was computed as in the original study by cropping all 674 ‘cell-type agnostic’ DNase profiles to a 10-bin (1,280 nt) region centred about the variant of interest, then summing across bins in the mean cropped profile.
- BPNet predicts nucleotide-resolution ChIP-nexus binding profiles for four TFs (Oct4, Sox2, Klf4 and Nanog) in mouse embryonic stem cells⁵. BPNet takes as input a DNA sequence of length 1,000 nt and outputs a 1,000-valued positive (+) and negative (−) strand profile for each of the four TFs (eight profiles in total per input). The parameters of BPNet were retrained as specified in the original release⁶⁴, and the resulting model was confirmed to recapitulate the published model BPNet-OSKN⁶⁵ by a visual inspection of attribution maps. Analyses for different prediction tasks used different profiles (Supplementary Table 1). In all BPNet analyses except those in Extended Data Fig. 8, y was computed using the profile contribution score defined in the original paper⁵. In the BPNet analysis for Extended Data Fig. 8, y was instead computed using PCA as described above and in Supplementary Fig. 3.
- The baseline CNN predicts ChIP-seq peaks for the human TF GABPA in GM12878 cells. This model takes as input a DNA sequence of length 200 nt and outputs a single probability. In our analysis of the effects of benign overfitting, y was computed as the logit of the output probability. The baseline CNN has not been previously published. Briefly, this model was trained to distinguish binary-labelled sequences: ChIP-seq peaks of GABPA in human GM12878 cells (positive labels) and DNase-seq peaks of GM12878 cells that did not overlap with any GABPA peaks (negative labels). Data were acquired from <https://zenodo.org/record/7011631> (data/GABPA_200.h5); 11,022, 1,574 and 3,150 sequences in the training, validation and test sets were respectively used. The CNN takes as input a DNA sequence of length 200 nt and outputs a probability. The hidden layers consist of three convolutional blocks each with max pooling (size 4 with stride 4), followed by a fully connected hidden layer with 128 units and an output layer to a single node with sigmoid activations. The number of filters and the kernel sizes from the first convolutional layer to the third are (32, 15), (64, 5) and (96, 5). The same padding was used for each convolutional layer. All hidden layers used ReLU activations and dropout⁶⁶, with rates given in order by 0.1, 0.2, 0.3 and 0.5. The CNN was trained to minimize the binary cross-entropy loss function using Adam⁶⁷, with default parameters and a batch size of 64. Early stopping was implemented to save model parameters at the epoch that corresponded to the largest area under the receiver-operator characteristic curve (AUROC) on the validation set, yielding the pre-overfitting DNN parameters. From that point, the CNN was trained for an additional 200 epochs, yielding the post-overfitting DNN parameters.

Attribution methods

Our analyses used attribution maps computed using a variety of methods, implemented as follows.

- ISM scores were computed by evaluating the scalar DNN prediction for every SNV of the sequence of interest.
- Saliency Maps scores were computed by evaluating the gradient of the scalar DNN prediction at the sequence of interest with respect to the one-hot encoding of that sequence.
- DeepSHAP scores were computed using the algorithm implemented in the DeepSTARR repository³⁴.
- DeepLIFT scores were computed using the algorithm implemented in the BPNet repository⁵.
- SmoothGrad scores were computed by averaging Saliency Maps over 50 noisy encodings of the sequence of interest. Each noisy encoding was computed by adding Gaussian noise (mean zero, standard deviation 0.25) to each of the $4L$ matrix elements of the one-hot encoding for the sequence of interest.

Standardization of attribution maps

Prior to comparing attribution maps, we standardized these maps to remove non-identifiable and/or non-meaningful degrees of freedom. Using $v_{l:c}$ to denote the attribution map value for character c at position l , this standardization was carried out as follows.

- For plotting sequence logos and computing attribution variation values: attribution map values were standardized using the transformation $v_{l:c} \rightarrow (v_{l:c} - \bar{v}_l)/\sigma$, where $\bar{v}_l = \frac{1}{4} \sum_{c'} v_{l:c'}$ and $\sigma^2 = \frac{1}{4L} \sum_{c,l} (v_{l:c} - \bar{v}_l)^2$. This transformation is essentially a gradient correction at each position⁶⁸, followed by a normalization with the square root of the total variation of the attribution scores across the sequence.
- For variant-effect analysis (Table 1): attribution map values were standardized using the transformation $v_{l:c} \rightarrow v_{l:c} - v_{l:s^{wt}}$, where s^{wt} is the wild-type sequence and s_l^{wt} is the character at position l in this sequence.
- For plotting heatmaps of additive and pairwise-interaction model parameters (Fig. 5c–e and Extended Data Fig. 8d): parameters were standardized as in ref. 26 using the ‘empirical gauge’.

Attribution variation computations

Attribution variation values were computed as follows.

- For consensus TF binding sites (Figs. 2 and 3), we first located all instances of the consensus TF binding site in the genome; the consensus sites used for each TF are listed in Supplementary Table 1. Using a baseline attribution method (Saliency Maps for ResidualBind-32, ISM for DeepSTARR and BPNet), we then ranked and manually pruned these genomic sequences to identify $m = 50$ putative functional and spatially isolated TF binding sites. For each of the m genomic sequences, we then computed an attribution map spanning the putative TF binding site plus 100 nucleotides on either side (that is, flanks). We then cropped these m attribution maps to span n_f nucleotides on either side of the consensus TF binding site. For Fig. 2 we used $n_f = 15$; for Fig. 3 we used $n_f = 0, 5, 10, 20, 30, \dots, 100$. For each of the m cropped attribution maps, the attribution variation was defined to be the Euclidean distance between the cropped attribution map and the average of the m cropped attribution maps.
- For weak TF binding sites (Extended Data Fig. 5), we first located all instances of variants of the consensus TF binding site in the genome having up to 2 naturally occurring mutations. For each group of putative binding sites having zero, one or two mutations, we then ranked and manually pruned these genomic

sequences to identify $m = 50$ putative functional and spatially isolated TF binding sites as above. For each of the $3m$ genomic sequences, we then computed an attribution map spanning the putative TF binding site plus 100 nucleotides on either side. We then cropped these $3m$ attribution maps to span $n_f = 50$ nucleotides on either side of the consensus TF binding site. For each of the $3m$ cropped attribution maps, the attribution variation was defined to be the Euclidean distance between the cropped attribution map and the average of the m cropped attribution maps from the ensemble of binding sites having zero mutations.

Context-averaged epistatic interactions

To compute the context-averaged epistatic interactions shown in Fig. 5c, first, we located all pairs of consensus AP-1 binding sites in the genome separated by no more than 20 nucleotides. We then ranked and pruned these pairs to identify $m = 50$ genomic sequences for further analysis. For each of the m genomic sequences, we used SQUID to infer a pairwise-interaction model spanning the pair of putative sites and six nucleotides on either side. Parameter values were standardized by expressing them in the empirical gauge²⁶. Additive parameters for each site were then cropped and averaged over the m sequences. The intra-site and inter-site pairwise-interaction parameters were similarly cropped and averaged. This process is schematized in Fig. 5a.

Variant-effect analysis

The ability of different attribution methods to predict variant effects, quantified in Table 1, was benchmarked as follows.

- The variant-effect data used to perform these benchmark studies are MPRA data from ref. 47. These data comprise measurements for the effects of nearly all SNVs in 15 disease-associated regulatory sequences, with each sequence ranging from 187 nt to 600 nt in length. For each assayed SNV, variant effect was quantified as the difference in measured activity between the variant and wild-type sequences.
- To compute attribution map predictions of variant effects, we extracted the genomic sequences centred about each of the 15 assayed regulatory sequences (2,048 nt per sequence for tests of ResidualBind-32 and Basenji-32; 393,216 nt per sequence for tests of Enformer). For each attribution method, the effect of an SNV, from character s_l^{wt} to character c at position l in each regulatory sequence s^{wt} , was quantified as $v_{l:c} - v_{l:s_l^{wt}}$. When inferring surrogate models using SQUID, we mutagenized the regulatory sequence and 400 nt of flanking sequence on either side.

Reporting summary

Further information on research design is available in the Nature Portfolio Reporting Summary linked to this article.

Data availability

The datasets, models and computational results used to support the findings in this paper are available on Zenodo at <https://doi.org/10.5281/zenodo.10047748> ref. 69. Datasets include the test set sequences held out during the training of ResidualBind-32, DeepSTARR and BPNet; the ChIP-seq peaks and background sequences used to train our three-layer CNN; and the CAGIS challenge dataset.

Code availability

SQUID is an open-source Python package based on TensorFlow⁷⁰. SQUID can be installed via pip (<https://pypi.org/project/squid-nn>) or GitHub (<https://github.com/evanseitz/squid-nn>). Documentation for SQUID is provided on ReadTheDocs (<https://squid-nn.readthedocs.io>). The code for performing all analyses in this paper is available on GitHub as well (<https://github.com/evanseitz/squid-manuscript>, ref. 71), and a static snapshot of this code is available on Zenodo⁶⁹.

References

1. Linder, J., Srivastava, D., Yuan, H., Agarwal, V. & Kelley, D. R. Predicting RNA-seq coverage from DNA sequence as a unifying model of gene regulation. Preprint at *bioRxiv* <https://doi.org/10.1101/2023.08.30.555582> (2023).
2. Avsec, Ž. et al. Effective gene expression prediction from sequence by integrating long-range interactions. *Nat. Methods* **18**, 1196–1203 (2021).
3. Dudnyk, K., Cai, D., Shi, C., Xu, J. & Zhou, J. Sequence basis of transcription initiation in the human genome. *Science* **384**, 694 (2024).
4. Jaganathan, K. et al. Predicting splicing from primary sequence with deep learning. *Cell* **176**, 535–548.e24 (2019).
5. Avsec, Ž. et al. Base-resolution models of transcription-factor binding reveal soft motif syntax. *Nat. Genet.* **53**, 354–366 (2021).
6. Chen, K. M., Wong, A. K., Troyanskaya, O. G. & Zhou, J. A sequence-based global map of regulatory activity for deciphering human genetics. *Nat. Genet.* **54**, 940–949 (2022).
7. Zhou, J. Sequence-based modeling of three-dimensional genome architecture from kilobase to chromosome scale. *Nat. Genet.* **54**, 725–734 (2022).
8. Koo, P. K. & Ploenzke, M. Deep learning for inferring transcription factor binding sites. *Curr. Opin. Syst. Biol.* **19**, 16–23 (2020).
9. Novakovsky, G., Dexter, N., Libbrecht, M. W., Wasserman, W. W. & Mostafavi, S. Obtaining genetics insights from deep learning via explainable artificial intelligence. *Nat. Rev. Genet.* **24**, 125–137 (2022).
10. Han, T., Srinivas, S. & Lakkaraju, H. Which explanation should I choose? A function approximation perspective to characterizing post hoc explanations. Preprint at <https://arxiv.org/abs/2206.01254> (2022).
11. Hooker, S., Erhan, D., Kindermans, P.-J. & Kim, B. A benchmark for interpretability methods in deep neural networks. In *Advances in Neural Information Processing Systems* Vol. 32 (2019).
12. Ancona, M., Ceolini, E., Öztireli, C. & Gross, M. Towards better understanding of gradient-based attribution methods for deep neural networks. Preprint at <https://arxiv.org/abs/1711.06104> (2017).
13. Simonyan, K., Vedaldi, A. & Zisserman, A. Deep inside convolutional networks: visualising image classification models and saliency maps. In *Workshop at International Conference on Learning Representations* (2014).
14. Shrikumar, A., Greenside, P. & Kundaje, A. Learning important features through propagating activation differences. In *Proc. 34th International Conference on Machine Learning* Vol. 70, ICML'17, 3145–3153 (JMLR.org, 2017).
15. Zhou, J. & Troyanskaya, O. G. Predicting effects of noncoding variants with deep learning-based sequence model. *Nat. Methods* **12**, 931–934 (2015).
16. Smilkov, D., Thorat, N., Kim, B., Viégas, F. & Wattenberg, M. SmoothGrad: Removing noise by adding noise. Preprint at <https://arxiv.org/abs/1706.03825> (2017).
17. Sundararajan, M., Taly, A. & Yan, Q. Axiomatic attribution for deep networks. Preprint at <https://arxiv.org/abs/1703.01365> (2017).
18. Lundberg, S. M. & Lee, S.-I. A Unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems* Vol. 30, 4768–4777 (Curran Associates, 2017).
19. Starr, T. N. & Thornton, J. W. Epistasis in protein evolution. *Protein Sci.* **25**, 1204–1218 (2016).
20. Weinreich, D. M., Lan, Y., Wylie, C. S. & Heckendorn, R. B. Should evolutionary geneticists worry about higher-order epistasis? *Curr. Opin. Genet. Dev.* **23**, 700–707 (2013).
21. Aghazadeh, A. et al. Epistatic net allows the sparse spectral regularization of deep neural networks for inferring fitness functions. *Nat. Commun.* **12**, 5225 (2021).
22. Zhou, J. et al. Higher-order epistasis and phenotypic prediction. *Proc. Natl Acad. Sci. USA* **119**, e2204233119 (2022).
23. Domingo, J., Baeza-Centurion, P. & Lehner, B. The causes and consequences of genetic interactions (epistasis). *Annu. Rev. Genomics Hum. Genet.* **20**, 433–460 (2019).
24. Otwinowski, J., McCandlish, D. M. & Plotkin, J. B. Inferring the shape of global epistasis. *Proc. Natl Acad. Sci. USA* **115**, E7550–E7558 (2018).
25. Poelwijk, F. J., Krishna, V. & Ranganathan, R. The context-dependence of mutations: a linkage of formalisms. *PLOS Comput. Biol.* **12**, e1004771 (2016).
26. Tareen, A. et al. MAVe-NN: learning genotype–phenotype maps from multiplex assays of variant effect. *Genome Biol.* **23**, 98 (2022).
27. Tonner, P. D., Pressman, A. & Ross, D. Interpretable modeling of genotype–phenotype landscapes with state-of-the-art predictive power. *Proc. Natl Acad. Sci. USA* **119**, e2114021119 (2022).
28. Ribeiro, M. T., Singh, S. & Guestrin, C. “Why should I trust you?”: explaining the predictions of any classifier. In *Proc. 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining* 1135–1144 (ACM, 2016).
29. Kinney, J. B., Murugan, A., Callan Jr, C. G. & Cox, E. C. Using deep sequencing to characterize the biophysical mechanism of a transcriptional regulatory sequence. *Proc. Natl Acad. Sci. USA* **107**, 9158–9163 (2010).
30. Jones, M. C. & Faddy, M. J. A skew extension of the t-distribution, with applications. *J. R. Stat. Soc. Ser. B* **65**, 159–174 (2003).
31. Tareen, A. & Kinney, J. B. Logomaker: beautiful sequence logos in Python. *Bioinformatics* **36**, 2272–2274 (2019).
32. Gordán, R. et al. Genomic regions flanking E-box binding sites influence DNA binding specificity of bHLH transcription factors through DNA shape. *Cell Rep.* **3**, 1093–1104 (2013).
33. Jolma, A. et al. DNA-dependent formation of transcription factor pairs alters their binding specificity. *Nature* **527**, 384–388 (2015).
34. de Almeida, B. P., Reiter, F., Pagani, M. & Stark, A. DeepSTARR predicts enhancer activity from DNA sequence and enables the de novo design of synthetic enhancers. *Nat. Genet.* **54**, 613–624 (2022).
35. Toneyan, S., Tang, Z. & Koo, P. Evaluating deep learning for predicting epigenomic profiles. *Nat. Mach. Intell.* **4**, 1088–1100 (2022).
36. Spitz, F. & Furlong, E. E. M. Transcription factors: from enhancer binding to developmental control. *Nat. Rev. Genet.* **13**, 613–626 (2012).
37. Bartlett, P. L., Long, P. M., Lugosi, G. & Tsigler, A. Benign overfitting in linear regression. *Proc. Natl Acad. Sci. USA* **117**, 30063–30070 (2020).
38. Chatterji, N. S. & Long, P. M. Finite-sample analysis of interpolating linear classifiers in the overparameterized regime. *J. Mach. Learn. Res.* **22**, 5721–5750 (2021).
39. Wang, Z. et al. Smoothed geometry for robust attribution. *Adv. Neural Inform. Process. Syst.* **33**, 13623–13634 (2020).
40. Alvarez-Melis, D. & Jaakkola, T. S. Towards robust interpretability with self-explaining neural networks. In *Proc. 32nd International Conference on Neural Information Processing Systems* 7786–7795 (Curran Associates Inc., 2018).
41. Majdandzic, A. et al. Selecting deep neural networks that yield consistent attribution-based interpretations for genomics. In *Machine Learning in Computational Biology* 131–149 (PMLR, 2022).
42. Zhang, C., Bengio, S., Hardt, M., Recht, B. & Vinyals, O. Understanding deep learning requires rethinking generalization. Preprint at <https://arxiv.org/abs/1611.03530> (2017).
43. Papagianni, A. et al. Capicua controls Toll/IL-1 signaling targets independently of RTK regulation. *Proc. Natl Acad. Sci. USA* **115**, 1807–1812 (2018).

44. Crocker, J. et al. Low affinity binding site clusters confer hox specificity and regulatory robustness. *Cell* **160**, 191–203 (2015).
45. Farley, E. K. et al. Suboptimization of developmental enhancers. *Science* **350**, 325–328 (2015).
46. Castro-Mondragon, J. A. et al. JASPAR 2022: the 9th release of the open-access database of transcription factor binding profiles. *Nucleic Acids Res.* **50**, D165–D173 (2021).
47. Kircher, M. et al. Saturation mutagenesis of twenty disease-associated regulatory elements at single base-pair resolution. *Nat. Commun.* **10**, 3583 (2019).
48. Shigaki, D. et al. Integration of multiple epigenomic marks improves prediction of variant impact in saturation mutagenesis reporter assay. *Hum. Mutat.* **40**, 1280–1291 (2019).
49. Kelley, D. R. et al. Sequential regulatory activity prediction across chromosomes with convolutional neural networks. *Genome Res.* **28**, 739–750 (2018).
50. Kim, S. & Wysocka, J. Deciphering the multi-scale, quantitative *cis*-regulatory code. *Mol. Cell* **83**, 373–392 (2023).
51. Georgakopoulos-Soares, I. et al. Transcription factor binding site orientation and order are major drivers of gene regulatory activity. *Nat. Commun.* **14**, 2333 (2023).
52. Koo, P. K., Majdandzic, A., Ploenzke, M., Anand, P. & Paul, S. B. Global importance analysis: an interpretability method to quantify importance of genomic features in deep neural networks. *PLoS Comput. Biol.* **17**, e1008925 (2021).
53. Weinreich, D. M., Lan, Y., Jaffe, J. & Heckendorf, R. B. The influence of higher-order epistasis on biological fitness landscape topography. *J. Stat. Phys.* **172**, 208–225 (2018).
54. Ackers, G. K., Johnson, A. D. & Shea, M. A. Quantitative model for gene regulation by lambda phage repressor. *Proc. Natl Acad. Sci. USA* **79**, 1129–1133 (1982).
55. Bintu, L. et al. Transcriptional regulation by the numbers: models. *Curr. Opin. Genet. Dev.* **15**, 116–124 (2005).
56. Segal, E. & Widom, J. From DNA sequence to transcriptional behaviour: a quantitative approach. *Nat. Rev. Genet.* **10**, 443–456 (2009).
57. Sherman, M. S. & Cohen, B. A. Thermodynamic state ensemble models of *cis*-regulation. *PLoS Comput. Biol.* **8**, e1002407 (2012).
58. Faure, A. J. et al. Mapping the energetic and allosteric landscapes of protein binding domains. *Nature* **604**, 175–183 (2022).
59. Tareen, A. & Kinney, J. B. Biophysical models of *cis*-regulation as interpretable neural networks. In *14th Conference on Machine Learning in Computational Biology (MLCB 2019)*; <https://doi.org/10.1101/835942>
60. Estrada, J., Wong, F., DePace, A. & Gunawardena, J. Information integration and energy expenditure in gene regulation. *Cell* **166**, 234–244 (2016).
61. Scholes, C., DePace, A. H. & Sánchez, Á. Combinatorial gene regulation through kinetic control of the transcription cycle. *Cell Syst.* **4**, 97–108.e9 (2017).
62. Park, J. et al. Dissecting the sharp response of a canonical developmental enhancer reveals multiple sources of cooperativity. *eLife* **8**, e41266 (2019).
63. Pedregosa, F. et al. Scikit-learn: machine learning in Python. *J. Mach. Learn. Res.* **12**, 2825–2830 (2011).
64. Avsec, Z. & Weilert, M. kundajelab/bpnet-manuscript: Publication release of BPNet manuscript code. Zenodo <https://zenodo.org/records/4294814> (2020).
65. Avsec, Z. et al. The Kipoi repository accelerates community exchange and reuse of predictive models for genomics. *Nat. Biotechnol.* <https://doi.org/10.1038/s41587-019-0140-0> (2019).
66. Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I. & Salakhutdinov, R. Dropout: a simple way to prevent neural networks from overfitting. *J. Mach. Learn. Res.* **15**, 1929–1958 (2014).
67. Kingma, D. P. & Ba, J. Adam: a method for stochastic optimization. Preprint at <https://arxiv.org/abs/1412.6980> (2014).
68. Majdandzic, A., Rajesh, C. & Koo, P. K. Correcting gradient-based interpretations of deep neural networks for genomics. *Genome Biol.* **24**, 1–13 (2023).
69. Seitz, E. evanseitz/squid-manuscript: SQUID manuscript workflow with outputs. Zenodo <https://doi.org/10.5281/zenodo.10047747> (2023).
70. Abadi, M. et al. TensorFlow: large-scale machine learning on heterogeneous systems. Software available from tensorflow.org (2015).
71. Seitz, E. & Koo, P. evanseitz/squid-nn: surrogate quantitative interpretability for deepnets. Zenodo <https://doi.org/10.5281/zenodo.11060672> (2023).

Acknowledgements

We thank Z. Tang, S. Toneyan, M. Kooshbaghi, C. Rajesh, J. Kaczmarzyk and C. Martí-Gómez for helpful discussions. This work was supported in part by: the Simons Center for Quantitative Biology at Cold Spring Harbor Laboratory; NIH grants R01HG012131 (P.K.K., E.E.S., J.B.K. and D.M.M.), R01HG011787 (J.B.K., E.E.S. and D.M.M.), R01GM149921 (P.K.K.), R35GM133777 (J.B.K.) and R35GM133613 (D.M.M.); and an Alfred P. Sloan Foundation Research Fellowship (D.M.M.). Computations were performed using equipment supported by NIH grant S10OD028632.

Author contributions

E.E.S., D.M.M., J.B.K. and P.K.K. conceived of the study. E.E.S. wrote the software and performed the analysis. E.E.S. designed the analysis with help from D.M.M., J.B.K. and P.K.K. J.B.K. and P.K.K. supervised the study. E.E.S., D.M.M., J.B.K. and P.K.K. wrote the paper.

Competing interests

The authors declare no competing interests.

Additional information

Extended data is available for this paper at <https://doi.org/10.1038/s42256-024-00851-5>.

Supplementary information The online version contains supplementary material available at <https://doi.org/10.1038/s42256-024-00851-5>.

Correspondence and requests for materials should be addressed to Justin B. Kinney or Peter K. Koo.

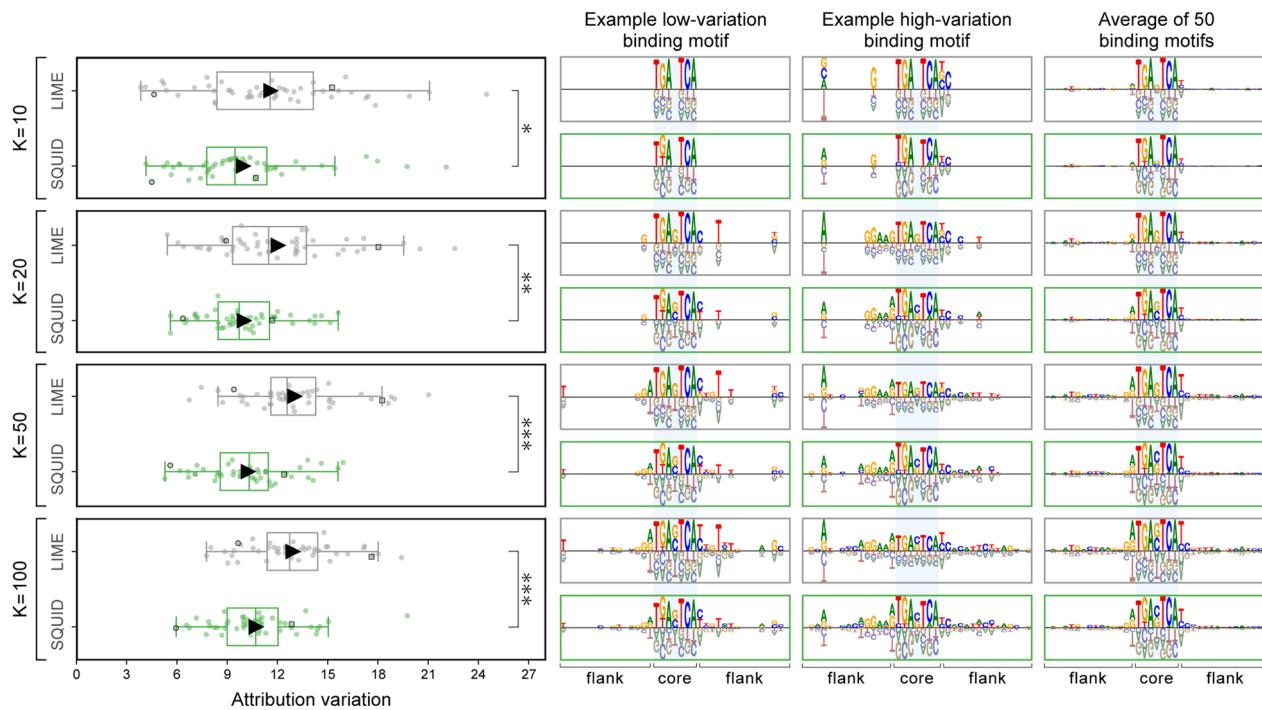
Peer review information *Nature Machine Intelligence* thanks the anonymous reviewers for their contribution to the peer review of this work.

Reprints and permissions information is available at www.nature.com/reprints.

Publisher's note Springer Nature remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.

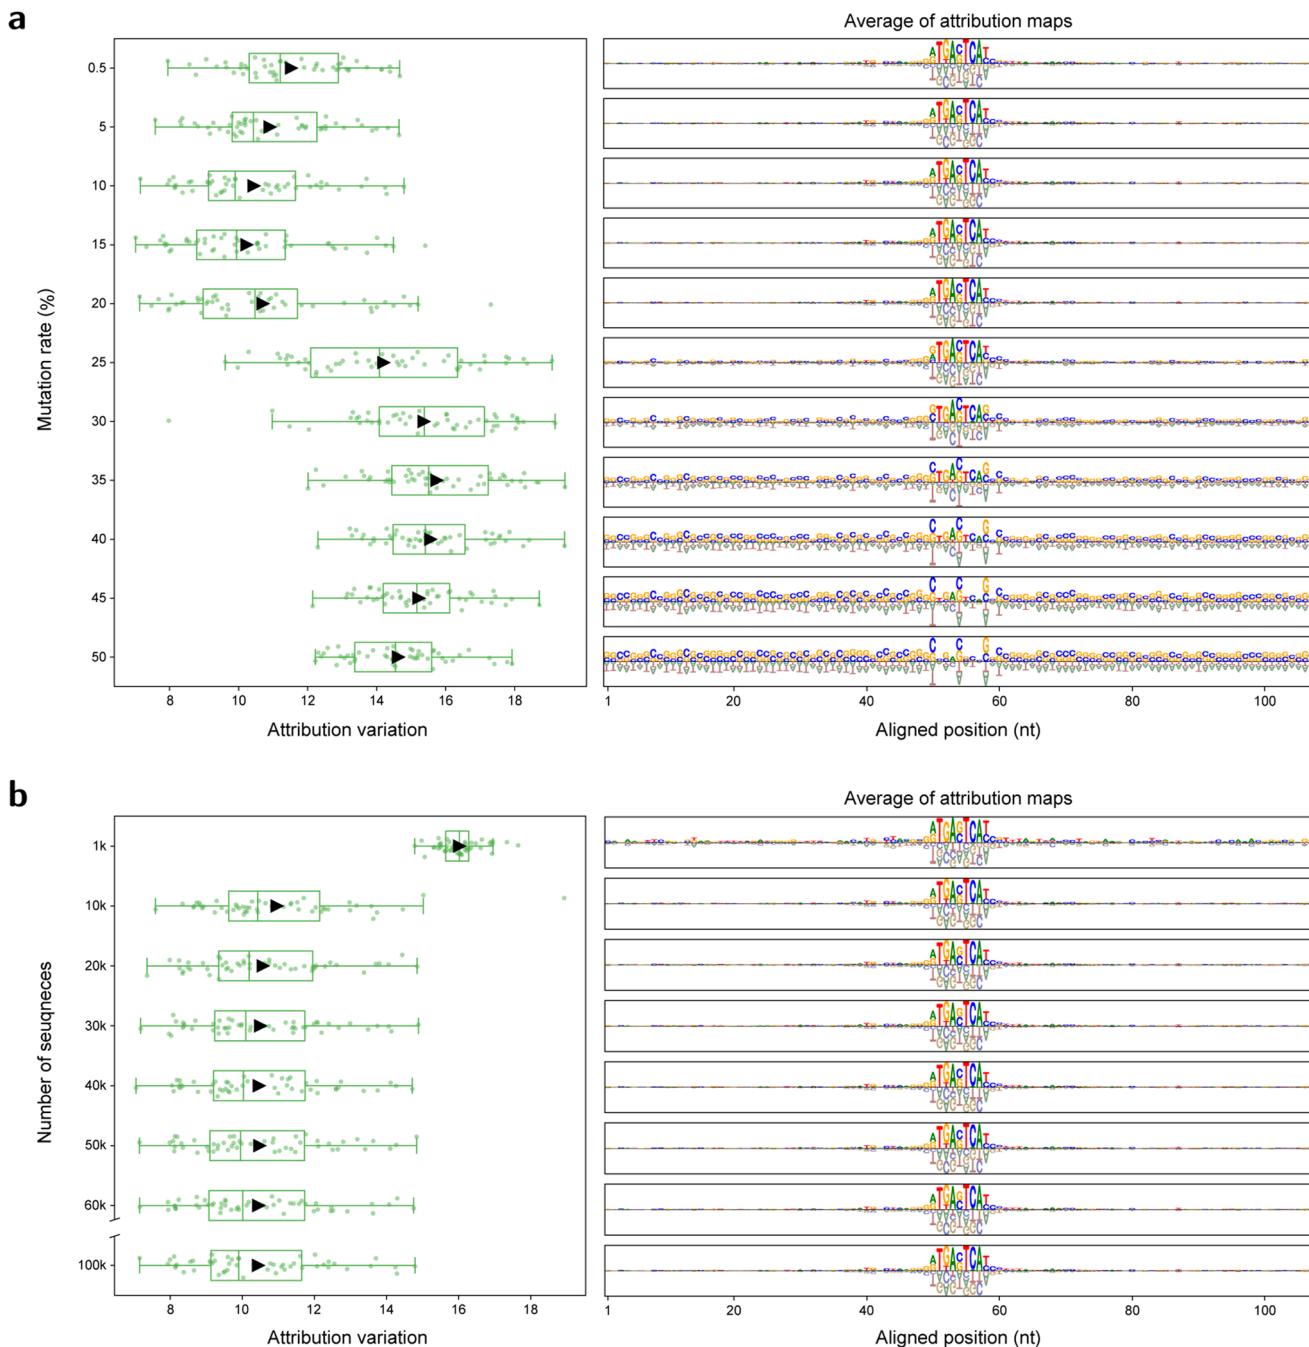
Springer Nature or its licensor (e.g. a society or other partner) holds exclusive rights to this article under a publishing agreement with the author(s) or other rightsholder(s); author self-archiving of the accepted manuscript version of this article is solely governed by the terms of such publishing agreement and applicable law.

© The Author(s), under exclusive licence to Springer Nature Limited 2024


Extended Data Fig. 1 | Additional comparisons of K-Lasso LIME to SQUID.

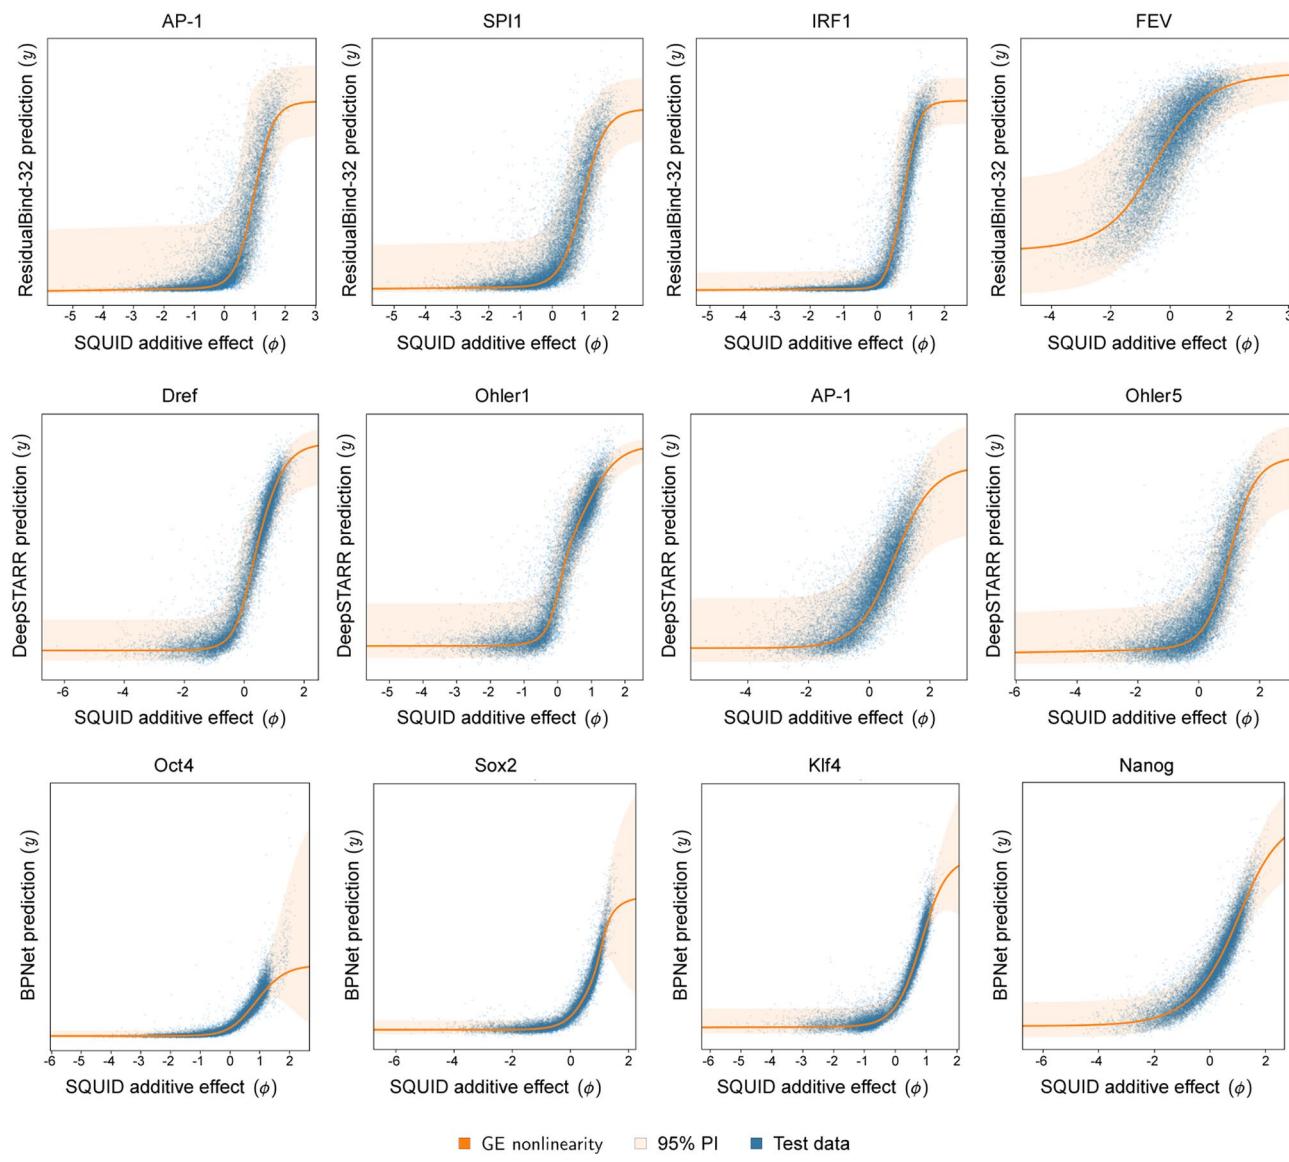
Shown are the results of analyses, performed as in Fig. 2b for $n = 50$ genomic sequences, comparing the performance of SQUID to the performance of the K-Lasso implementation of LIME for four different values of K . P values were computed using a one-sided Mann-Whitney U test; ***, $p < 0.001$. We note that the attribution variation values obtained for SQUID in these tests varied systematically with the choice of K . The reason is as follows. The K-Lasso LIME algorithm produces sparse attribution maps that have only K nonzero

parameters. Consequently, the variation observed in K-Lasso LIME attribution maps systematically decreases as K decreases. This gives K-Lasso LIME an unfair advantage in the attribution variation test described in Main Text and in Methods. To fairly compare K-Lasso LIME to SQUID in this figure, we therefore modified this test. In the analysis of each *in silico* MAVE, the attribution map elements inferred by SQUID were first set to zero at the same positions where all K-Lasso LIME attribution map elements were exactly zero. Attribution variation values were then calculated as described in Main Text and in Methods.



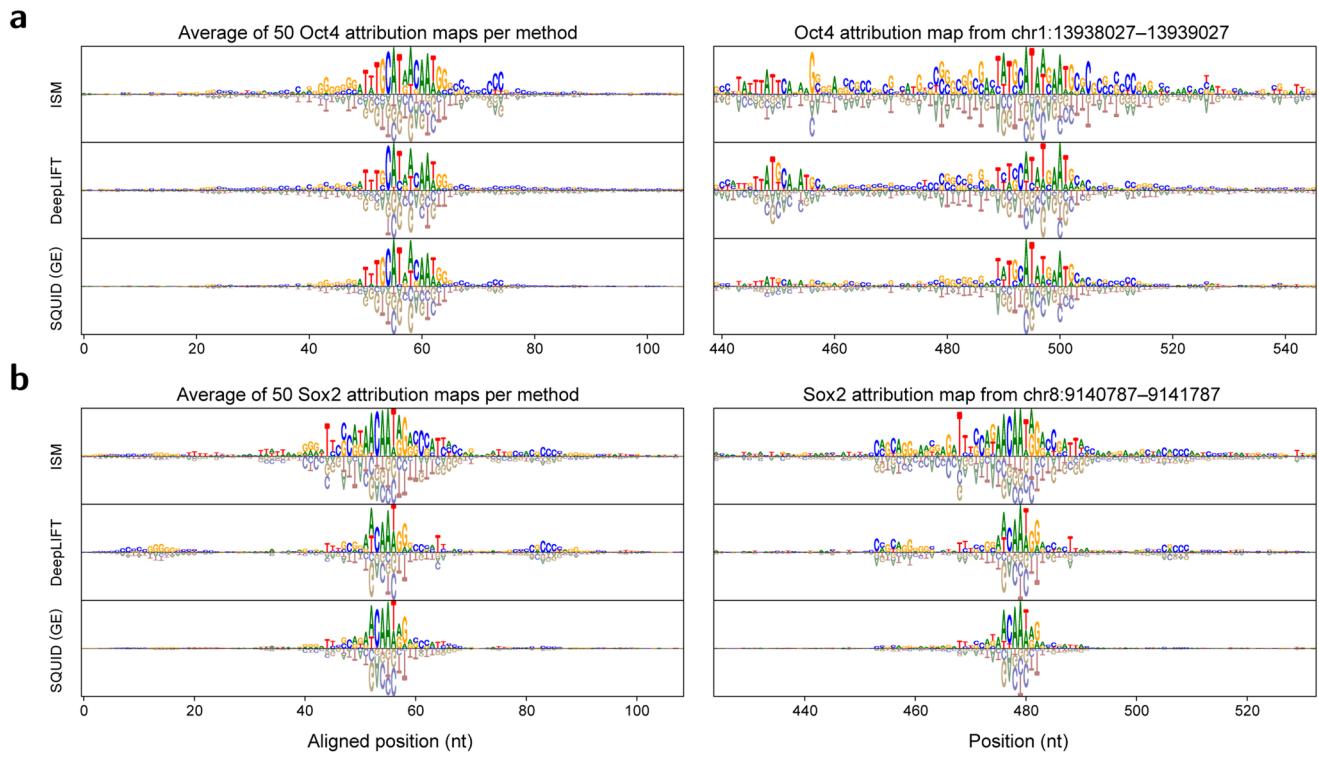
Extended Data Fig. 2 | Influence of mutation rate and library size on SQUID attribution maps. **a,b**, attribution variation (left) and average attribution maps (right) found for 50 SQUID attribution maps computed for the TF AP-1 as in Fig. 2 but using *in silico* MAVE libraries having (a) variable mutation rate (r) and

fixed number of mutagenized sequences $N = 100,000$, or (b) fixed mutation rate $r = 10\%$ and variable number of mutagenized sequences (N). All SQUID attribution maps were computed using additive models with GE nonlinearities, followed by cropping these maps using a flank size of $n_f = 50$ nt.

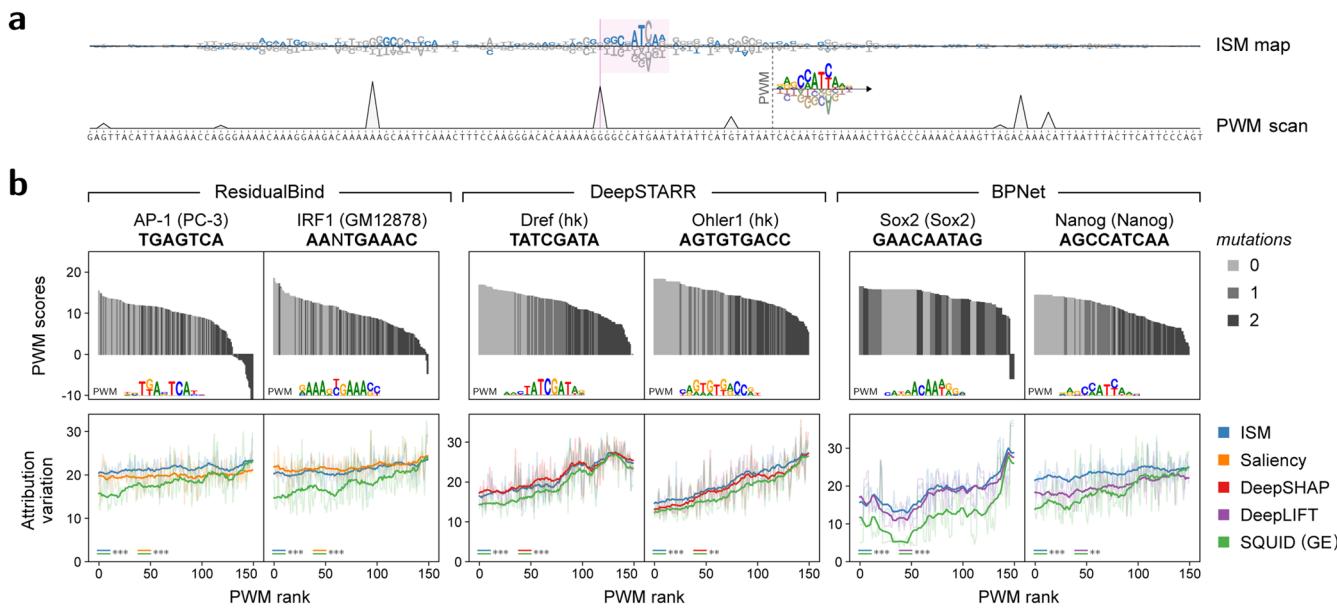


■ GE nonlinearity □ 95% PI ■ Test data

Extended Data Fig. 3 | Nonlinearities and noise across DNNs and TFs. Examples of the GE nonlinearities and heteroscedastic noise models inferred in the SQUID (GE) analyses performed for Fig. 3. Each plot shows results for a representative sequence among the $n = 50$ sequences analyzed for each combination of DNN and TF.

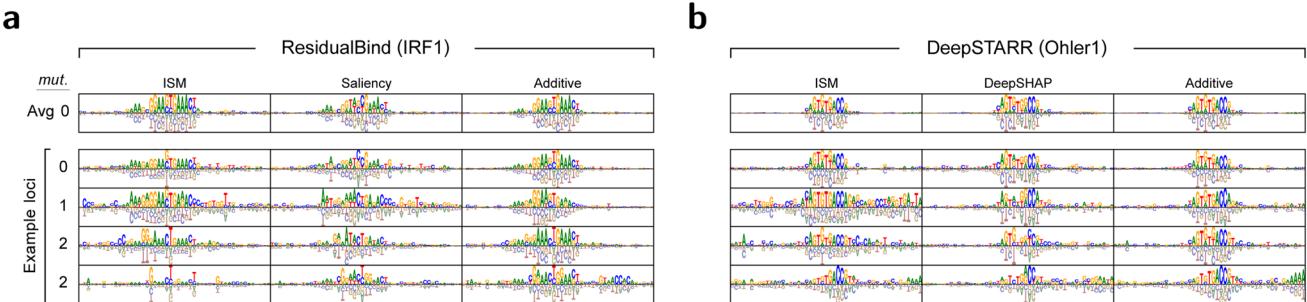


Extended Data Fig. 4 | Average and example binding motifs for Oct4 and Sox2. **a**, Oct4 motifs, centered on the putative binding site TTTGCAT. **b**, Sox2 motifs, centered on the putative binding site GAACAATAG. TF binding motifs are from attribution maps computed for BPNet and plotted as in Fig. 3c.



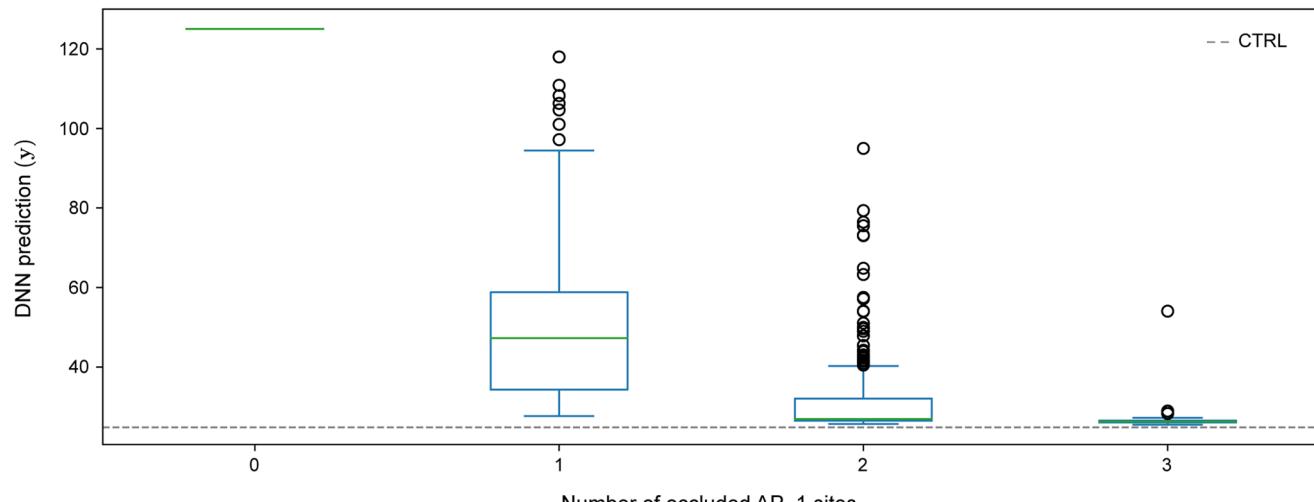
Extended Data Fig. 5 | Benchmark analysis of attribution variation for putative weak TF binding sites. a. a, ISM attribution map given by BPNet for a representative genomic sequence containing multiple putative weak binding sites for the mouse TF Nanog (top). Blue represents the wild-type nucleotides while gray represents other nucleotides. PWM scores (given by PWM for Nanog motif) are displayed across the genomic sequence (bottom). Only positive PWM scores are shown. **b,** For each TF and DNN, plots show attribution variation values

for 150 putative TF binding sites plotted against putative binding site strength. Bold lines indicate signals smoothed with a sliding window of 20 nt. Stars indicate P values computed using the one-sided Mann-Whitney U test: **, $0.001 \leq p < 0.01$; ***, $p < 0.001$. PWM scores for each of the 150 putative sites are shown above, along with a logo representation of the PWM used. Each site is represented by a gray bar shaded according to the number of mutations (0, 1, or 2) in the core of the putative site.



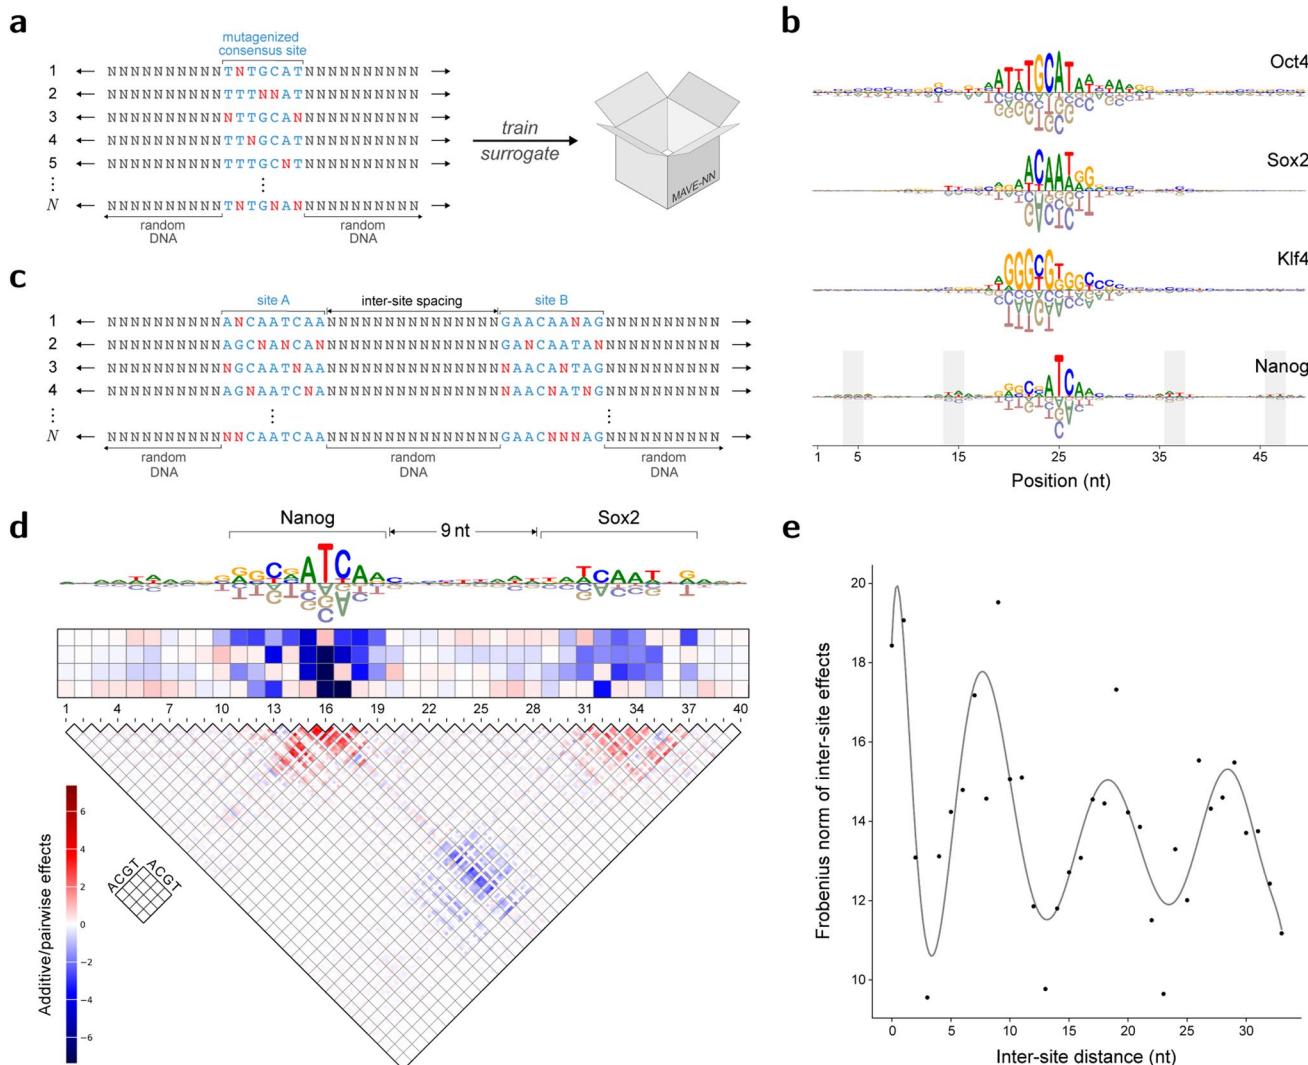
Extended Data Fig. 6 | Attribution maps computed for strong and weak TF binding sites. **a,b,** Top row shows the average of 50 attribution maps in the 0-mutation ensemble computed for IRF1 using ResidualBind-32 (**a**), and for

Ohler using DeepSTARR (**b**). Remaining rows show attribution maps for four representative genomic loci with the central putative binding site having varying numbers of mutations from the consensus binding site.

**Extended Data Fig. 7 | Occlusion analysis of AP-1 binding site effects.**

Occlusion study based on the wild-type sequence investigated in Figures 6d and 6e. Occlusions were performed for every combination of one, two, or three occluded motifs, with the DNN prediction taken independently for each instance ($n=100$ for each boxplot). In each occluded sequence, the corresponding AP-1 core (7-mer) sites were scrambled using a uniform probability of nucleotides

at each position. The baseline score (CTRL) was calculated from the median of predictions corresponding to $n=100$ instances of a dinucleotide shuffle over the full (2048 nt) sequence. The DNN prediction rapidly approaches the genomic baseline as additional binding sites are occluded. Boxplot lines represent median, upper quartile, and lower quartile; whiskers represent $1.5 \times$ the interquartile range.



Extended Data Fig. 8 | SQUID supports global DNN interpretations. **a**, In silico MAVE libraries used by SQUID to infer surrogate models of global TF specificity. Each sequence in the library contains a partially-mutagenized version of a consensus TF binding site embedded within random DNA. **b**, Additive surrogate models for BPNet inferred by SQUID using libraries of the form in panel **a** for the mouse TFs Oct4, Sox2, Klf4, and Nanog. Gray bars indicate 10.5 nt periodicity on either side of the inferred Nanog motif. **c**, Format of the in silico MAVE libraries used to study global TF-TF epistatic interactions. Each sequence in the library contains two partially mutagenized consensus TF binding sites embedded a fixed distance apart (0 nt to 32 nt) within random DNA. **d**, Pairwise surrogate model inferred for BPNet by SQUID using a library of the form in **c** and putative binding

sites for Nanog and Sox2. **e**, Distance-dependence of inter-site interactions between Nanog and Sox2. Dots show the Frobenius norm of inter-site pairwise interactions, inferred as in **d**, using libraries having different distances between the embedded Nanog and Sox2 sites. The occurrence of periodic inter-site interaction minima occurred at distances where the central nucleotides in Sox2 (that is, AA) overlapped with the periodic Nanog-associated flanking signals. Black line is from a least-squares fit of a polynomial of degree 10. **d,e**, Libraries comprised of 500,000 sequences. BPNet profiles were projected to scalars using PCA, which yielded substantially lower background noise than obtained using profile contribution scores.

Extended Data Table 1 | SQUID computation times

DNN	GP map	Sequence length (nt)	GP map parameters	Mutagenizer (minutes)	Predictor (minutes)	Training (minutes)	Training (epochs)	Total time (minutes)
DeepSTARR	Additive	250	1000	0.18±0.01	0.04±0.00	13.3±2.9	180.9±39.0	13.5±2.9
DeepSTARR	Pairwise	30	6960	0.09±0.00	0.05±0.01	4.5±2.8	102.1±63.4	4.6±2.8
ResidualBind-32	Additive	2048	8192	0.78±0.00	2.19±0.01	4.8±1.1	61.3±15.6	7.8±1.1
ResidualBind-32	Pairwise	30	6960	0.12±0.00	2.19±0.01	3.2±0.2	41.6±3.5	5.5±0.2
BPNet	Additive	1000	4000	0.34±0.00	17.36±0.01	11.3±2.7	138.7±34.0	29.0±2.7
BPNet	Pairwise	30	6960	0.09±0.00	17.35±0.01	4.8±2.4	62.6±33.2	22.2±2.4
Enformer	Additive	768	3072	0.37	1540.80	4.7±1.1	82.9±20.0	1545.9±1.1
Enformer	Pairwise	30	6960	0.09	1540.80	1.8±0.1	33.8±2.4	1542.7±0.1

Shown are SQUID computation times for the different stages of analysis of various genomic DNNs using various types of surrogate models. All analyses were performed using 100,000 variant sequences mutagenized at 10% per nucleotide, a GE measurement process with a skewed-T noise model, and early stopping. Error bars indicate standard deviations over 10 analysis runs. For DeepSTARR, ResidualBind-32, and BPNet, analyses were performed for 10 different *in silico* MAVE datasets. For Enformer, 10 analyses were performed on the same *in silico* MAVE dataset due to each forward-pass being slower for Enformer than for the other DNNs. Shorter sequences were used when inferring pairwise-interaction models to keep the number of pairwise-interaction model parameters comparable to the number of additive model parameters. DeepSTAR and ResidualBind-32 analyses were performed using a V100 GPU in Google Colab. BPNet and Enformer analyses were performed using an RTX A4000 GPU on a private server.

Extended Data Table 2 | GE modelling is robust to the number of hidden nodes

DNN	Measurement Process	GE Hidden Nodes	R^2	Epochs	test I_pred (bits)	max I_var (bits)	max val_I_var (bits)
DeepSTARR	Linear	N/A	0.70±0.03	98.4±23.4	1.2±0.1	0.5±0.1	0.5±0.1
DeepSTARR	GE	1	0.86±0.03	112.4±24.7	1.3±0.1	1.2±0.1	1.3±0.3
DeepSTARR	GE	50	0.86±0.03	202.1±41.5	1.3±0.1	1.3±0.1	1.2±0.1
DeepSTARR	GE	1000	0.86±0.03	146.8±54.7	1.3±0.1	1.3±0.1	1.2±0.1
ResidualBind-32	Linear	N/A	0.70±0.18	113.8±37.8	1.4±0.2	0.2±0.4	0.0±0.4
ResidualBind-32	GE	1	0.93±0.04	136.9±79.4	2.3±0.2	2.4±0.2	2.0±0.2
ResidualBind-32	GE	50	0.94±0.04	116.7±80.3	2.3±0.2	2.4±0.2	2.0±0.3
ResidualBind-32	GE	1000	0.93±0.05	127.0±47.8	2.3±0.2	2.4±0.2	2.1±0.2

Shown are results for SQUID analyses performed using surrogate models having either no GE nonlinearity (Linear) or a GE nonlinearity (GE) computed using different numbers of hidden nodes (that is, tanh sigmoidal components; see Methods). Analyses were performed as described in Extended Data Table 1, except that in each test the 10 *in silico* MAVEs were generated for the same genomic locus (as opposed to different genomic loci). For DeepSTARR, an additive GP-map covering 250 nt was inferred. For ResidualBind-32, an additive GP-map covering 2048 nt was inferred. Error bars represent standard deviations over the 10 *in silico* MAVE datasets.

Reporting Summary

Nature Portfolio wishes to improve the reproducibility of the work that we publish. This form provides structure for consistency and transparency in reporting. For further information on Nature Portfolio policies, see our [Editorial Policies](#) and the [Editorial Policy Checklist](#).

Statistics

For all statistical analyses, confirm that the following items are present in the figure legend, table legend, main text, or Methods section.

n/a Confirmed

- The exact sample size (n) for each experimental group/condition, given as a discrete number and unit of measurement
- A statement on whether measurements were taken from distinct samples or whether the same sample was measured repeatedly
- The statistical test(s) used AND whether they are one- or two-sided
Only common tests should be described solely by name; describe more complex techniques in the Methods section.
- A description of all covariates tested
- A description of any assumptions or corrections, such as tests of normality and adjustment for multiple comparisons
- A full description of the statistical parameters including central tendency (e.g. means) or other basic estimates (e.g. regression coefficient) AND variation (e.g. standard deviation) or associated estimates of uncertainty (e.g. confidence intervals)
- For null hypothesis testing, the test statistic (e.g. F , t , r) with confidence intervals, effect sizes, degrees of freedom and P value noted
Give P values as exact values whenever suitable.
- For Bayesian analysis, information on the choice of priors and Markov chain Monte Carlo settings
- For hierarchical and complex designs, identification of the appropriate level for tests and full reporting of outcomes
- Estimates of effect sizes (e.g. Cohen's d , Pearson's r), indicating how they were calculated

Our web collection on [statistics for biologists](#) contains articles on many of the points above.

Software and code

Policy information about [availability of computer code](#)

Data collection doi.org/10.5281/zenodo.10047748

Data analysis

All used software and custom algorithms are listed in <https://github.com/evanseitz/squid-manuscript> with a mirror of this repository additionally including all models and outputs at doi.org/10.5281/zenodo.10047748. The front-facing version of the code can be found in <https://github.com/evanseitz/squid-nn>. Our analysis employed these previously-published DNNs: Basenji-32, BPNet, DeepSTARR, Enformer, and ResidualBind-32. The Basenji-32 and ResidualBind-32 trained models are from <https://doi.org/10.5281/zenodo.6464031>; the DeepSTARR trained model is from <https://doi.org/10.5281/zenodo.5502060>; the BPNet model is from <https://doi.org/10.5281/zenodo.3371163>; and the Enformer model is from <https://doi.org/10.5281/zenodo.5098375>. Surrogate modeling with MAVE-NN was performed using MAVE-NN 1.02. Surrogate modeling with LIME was implemented using scikit-learn 1.3.2.

For manuscripts utilizing custom algorithms or software that are central to the research but not yet described in published literature, software must be made available to editors and reviewers. We strongly encourage code deposition in a community repository (e.g. GitHub). See the Nature Portfolio [guidelines for submitting code & software](#) for further information.

Data

Policy information about [availability of data](#)

All manuscripts must include a [data availability statement](#). This statement should provide the following information, where applicable:

- Accession codes, unique identifiers, or web links for publicly available datasets
- A description of any restrictions on data availability
- For clinical datasets or third party data, please ensure that the statement adheres to our [policy](#)

The datasets, models, and computational results used to support the findings in this paper are available on Zenodo at <https://doi.org/10.5281/zenodo.10047748>. Datasets include the test set sequences held out during the training of ResidualBind-32, DeepSTARR, and BPNet; the ChIP-seq peaks and background sequences used to train our 3-layer CNN; and the CAGIS challenge dataset.

Human research participants

Policy information about [studies involving human research participants and Sex and Gender in Research](#).

Reporting on sex and gender

n/a

Population characteristics

n/a

Recruitment

n/a

Ethics oversight

n/a

Note that full information on the approval of the study protocol must also be provided in the manuscript.

Field-specific reporting

Please select the one below that is the best fit for your research. If you are not sure, read the appropriate sections before making your selection.

Life sciences

Behavioural & social sciences

Ecological, evolutionary & environmental sciences

For a reference copy of the document with all sections, see nature.com/documents/nr-reporting-summary-flat.pdf

Life sciences study design

All studies must disclose on these points even when the disclosure is negative.

Sample size

For our analysis of attribution errors, we first located all instances of the consensus TF binding site in the genome. Using a baseline attribution method, we then ranked and manually pruned these sequences to identify a set of putative functional and spatially-isolated TF binding sites. We found the chosen numbers to be well representative of the population mean. To determine the number of sequences used in constructing each MAVE dataset, we performed controlled experiments analyzing the effects of library size on attribution error. From this analysis, we determined that 50 attribution maps were sufficient to compare effect sizes. Sample size calculations were not relevant to our studies, which were purely computational.

Data exclusions

As described above, putative TF binding sites were pruned based on rank and spatial isolation.

Replication

We evaluated existing deep learning models which have fixed weights. Since the weights are fixed, direct replication is not relevant for this study, which was purely computational. Instead, we analyzed our results over different deep learning models.

Randomization

For all surrogate models in this study, the weights and biases were initialized randomly.

Blinding

Held-out test data was used to evaluate the performance of all surrogate models in this study. The models did not have access to the data in the held-out dataset, and this data was only used post-training during model evaluation. Blinding was not relevant to our studies, which were purely computational.

Reporting for specific materials, systems and methods

We require information from authors about some types of materials, experimental systems and methods used in many studies. Here, indicate whether each material, system or method listed is relevant to your study. If you are not sure if a list item applies to your research, read the appropriate section before selecting a response.

Materials & experimental systems

n/a	Involved in the study
<input checked="" type="checkbox"/>	Antibodies
<input checked="" type="checkbox"/>	Eukaryotic cell lines
<input checked="" type="checkbox"/>	Palaeontology and archaeology
<input checked="" type="checkbox"/>	Animals and other organisms
<input checked="" type="checkbox"/>	Clinical data
<input checked="" type="checkbox"/>	Dual use research of concern

Methods

n/a	Involved in the study
<input checked="" type="checkbox"/>	ChIP-seq
<input checked="" type="checkbox"/>	Flow cytometry
<input checked="" type="checkbox"/>	MRI-based neuroimaging